DC-Debates: Difficulty-Calibrated Debate Transcripts as Stimuli for Scalable Oversight Research with Human Participants

Gabriel Recchia[1], Chatrik Singh Mangat[2], Suhasini Satish Rao[3], Sanika Khadkikar[4]
[1]Modulo Research
[2]Vector Research
[3]University of Alberta
[4]The Pennsylvania State University
gabe@moduloresearch.com

Abstract

Approaches to scalable oversight involving human judges, such as AI safety via debate, will be more likely to achieve success if systematic errors in judgment can be minimized. Research on approaches to reducing this error, such as different ways of aggregating participant judgments or manipulating instructions to human participants, is more likely to generalise well if conducted using stimuli analogous to the kinds of stimuli that we expect strong trained AI debaters to generate. This requires well-designed transcripts that are difficult enough to leave room for interventions to improve judge accuracy, but tractable enough that an attentive judge has some hope of arriving at the correct answer. We describe the construction of 24 debate transcripts spanning 12 questions in three domains (astrophysics research, brainteasers, and reward-hacking evaluations) that assume an underlying recursive protocol, with section-level balancing across 11 stylistic and rhetorical dimensions. We report a preregistered between-subjects study with human judges (N = 1,282) which estimates per-transcript posterior difficulty on the logit scale and partitions the 24 transcripts into a four-tier ordinal difficulty matrix. Our results show that our stimuli span a range of difficulty levels, providing a calibrated stimulus pool for subsequent research. We also report several failure modes in which honest debaters did not bring judges to the correct answer despite faithful argumentation, including a recursive-protocol-specific failure in which the dishonest debater steers recursion onto the less important of two subclaims. All transcripts, judge instructions, and analysis code are released as supplementary materials to facilitate follow-up work.

Code, transcripts, and analysis — https://osf.io/8wm72/
Preregistration — https://osf.io/wd93z/

1 Introduction

Debate (Irving, Christiano, and Amodei 2018) is one of the most actively studied approaches for maintaining human oversight of AI systems in a manner that scales with their growing capabilities. In this setup, two instances of a strong but potentially untrusted system are placed in an adversarial debate over the best answer to a question of interest, and a less capable judge (e.g., a human or trusted model) decides between them. The hope is that telling the full truth provides a structural advantage that survives even as questions become too hard for the judge to settle without AI assistance. A growing body of work involving human judges explores whether this advantage survives in practice, often investigating how debate compares to 'consultancy' baselines involving only a single persuader (Michael et al. 2023; Khan et al. 2024). The most commonly used dataset in empirical debate studies is QuALITY (Pang et al. 2021). However, QuALITY is limited for this purpose in various ways: it only contains questions about fictional short stories which are hidden from judges, debaters' advantages over judges are primarily due to information asymmetry, and it is unclear how representative this domain is to those we expect oversight methods to be realistically deployed in, such as scientific discovery and AI alignment research. Furthermore, despite continued recognition of the importance of experiments involving human judgments for scalable oversight (Irving and Askell 2019; Cotra 2021; Anwar et al. 2024; Voudouris et al. 2026), QuALITY remains the only public dataset explicitly designed for facilitating debate research with human participants, and one of very few designed for facilitating scalable oversight research more broadly (see Related Work).

Irving and Askell (2019) outline various desiderata for stimuli likely to be most effective for debate research involving human participants, such as that they should involve questions for which true answers are known, false answers are plausible (at least to judges), and there are sufficient facts available that can be verified by the judge so that the judge has some realistic hope of determining the correct answer. For a false answer to be plausible, the debate must not hinge on a single checkable fact that the honest debater can simply reveal, yet it must not be so obscure that a judge has little hope of determining the correct answer in the span of the experiment. Additionally, it is best if debates are balanced with respect to surface criteria likely to bias human participants, such as the number of verifiable facts deployed by each debater, each debater's perceived level of confidence, professionalism, use of technical jargon, and so on, so that these do not add unnecessary variance. Producing transcripts that meet these criteria is a substantial undertaking. Once produced, their difficulty for any particular judge population is also difficult to predict from inspection. In our own pilot work, judges' assessments of which arguments were most persuasive were frequently surprising to us as transcript authors. Many of the most natural research questions in this area (e.g., if accuracy is raised by increasing the amount of time available to judges, or aggregating multiple judges, or changing instructions) require a stimulus pool that already spans a range of empirically calibrated difficulty levels. Ideally, the pool contains stimuli that are not so easy that they make it difficult to detect an intervention's effect, but not so hard that even a strong intervention would leave judges flipping a coin. A pool that includes stimuli with a range of difficulty levels also enables study of whether interventions will generalize across difficulty levels.

This paper describes a set of 24 debate transcripts built to meet this need. The transcripts are derived from 12 questions in three domains chosen to be challenging for general educated adults: astrophysics research controversies, physics/probability brainteasers, and judgments about whether reward hacking has occurred in LLM coding agent trajectories. For each question, we constructed two transcripts that share content but differ in argumentative choices. While the debates are fictional in the sense that they represent transcripts authored by hand to resemble plausible debater output rather than the results of actual debates between trained AI or human debaters, they follow a rule set inspired by Barnes and Christiano (2020), who describe a recursive protocol and their reasons for adopting it. We also conducted extensive section-level balancing of stylistic and rhetorical features across the two debaters within each transcript and across the four debate sections shown to judges.

Our contributions. Our main empirical contribution is a preregistered between-subjects study in which each of 1,282 Prolific participants read one transcript across four sequential parts, recording their credences in each debater's position after each part, as well as the reasons behind their judgments in free text. We fit a Bayesian hierarchical model with transcripts nested within questions nested within domains, derive a posterior over each transcript's total conditional difficulty, and partition the 24 transcripts into a four-tier difficulty classification with posterior uncertainty on tier membership. We also describe failure modes that we observed during construction and across two pilots and the main study. This includes a recursive-protocol-specific dynamic that we believe deserves more attention, whereby dishonest debaters can take advantage of judges' poor understandings of what points are important vs. unimportant, denying honest debaters the opportunity to provide the judges enough context to understand why their primary objections are more important than their opponents' weaker points. All code, transcripts, and analysis, including intermediate versions of transcripts used in the pilots and the qualitative study, have been made publicly available[1], along with a website containing an interactive visual summary of key findings[2].

2 Related Work

Empirical scalable oversight studies with human judges. Irving, Christiano, and Amodei (2018) and Irving and Askell (2019) presented early informal examples of human debate experiments intended to inspire future related work. Work by Barnes et al. led to innovations debate protocols can use to handle various dishonest strategies (Barnes and Christiano 2020), although also suggested conditions under which dishonest strategies can defeat honest ones (Barnes et al. 2020). Parrish et al. (2022a,b) found that single-turn or two-turn debate does not help humans on reading comprehension tasks based on QuALITY (Pang et al. 2021). However, Michael et al. (2023) found that debate provided a significant improvement to judge accuracy over consultancy, with the main source of error arising from honest debaters making suboptimal choices regarding what evidence to cite, followed by judge inattentiveness. Khan et al. (2024) found that in addition to the boost in accuracy, human judges are better calibrated in debate than in consultancy. Rein (2024) reported an inability to bridge a very large domain expertise gap in two pilot experiments, but attributed this failure to contingent features of the setup rather than a problem with their protocol. Rahman et al. (2025) reported that debate boosts human judge accuracy over consultancy for domains involving COVID-19 claims and climate change.

Recursive debate. Most empirical debate work employs linear debate protocols, but theoretical debate work (Irving, Christiano, and Amodei 2018; Brown-Cohen, Irving, and Piliouras 2025) places a strong emphasis on recursive protocols. Barnes and Christiano (2020) ran human experiments with recursive debate and introduced further components like cross-examination and meta-debate, but were unable to overcome the issue of "obfuscated arguments" (Barnes et al. 2020).

Stimuli for scalable oversight. Michael et al. (2023), Khan et al. (2024), and Rahman et al. (2025) have released debate transcripts from their experiments. The first two are restricted to the QuALITY dataset, whereas those of Rahman et al. (2025) concern controversial factuality claims on COVID-19 and climate change. FindTheFlaws (Recchia et al. 2026) can be used to construct debate transcripts and train critique models (Saunders et al. 2022). Kenton et al. (2024) test AI debate on a multitude of tasks beyond QuALITY. Elasky, Nakasako, and Goyal (2026) conduct AI debates based on the CodeContests+ (Wang et al. 2025) and ARC-AGI-2 (Chollet et al. 2025) datasets. Pallavi Sudhir, Kaunismaa, and Panickssery (2025) offer a benchmark intended to support evaluations of scalable oversight protocols, while Engels et al. (2025) suggest a general framework for such comparisons. Other datasets that may prove useful for empirical debate or verification research include Sabotage Bench (Gan et al. 2026), EvilGenie (Gabor, Lynch, and Rosenfeld 2025), ImpossibleBench (Zhong, Raghunathan, and Carlini 2025) and the conceptual reasoning dataset of Cooper et al. (2026).

3 Methods

3.1 Constructing the stimulus set

Question selection We selected 12 questions across three domains (four per domain). Questions in the "brainteasers" domain (Internal Temperature of Stars, Switch, Bergs, Probability) all had a counterintuitive aspect, such that a dishonest debater can attempt to leverage an intuitive but incorrect argument. Switch and Bergs are variations on *Think-*

[1] https://github.com/scalable-oversight/dc-debates/

[2] https://moduloresearch.com/dcdebates/

ing Physics questions (Epstein 2002) previously explored by Barnes et al. (2020), and Probability Puzzle is borrowed from Coelho (2024), which in turn restates a 'paradox' often attributed to Elchanan Mossel (Jin 2018). The coding questions (Rust, Prefix-sum, ORM library, Database deletion) ask whether an excerpt of a (sometimes substantially modified) agent transcript drawn from METR's MALT-public dataset (Neev Parikh 2025) contains reward hacking; the correct answer is "yes" in two of the four cases.

The astrophysics questions (Little Red Dots, Early Universe Galaxies, Large Linear Structure, Gravitational Lens Modelling) concern resolved or near-resolved controversies in astrophysics research. The debate in Gravitational Lens Modelling is over two competing explanations of the same astrophysical system (J1721+8842), where Honest defends the findings of Dux et al. (2025) while Dishonest defends an older, since disproven model of Mangat et al. (2021). We chose this case in part because co-author Mangat was familiar with the original model from his earlier work on this system. Similarly, Large Linear Structure pits an Honest debater's claim, confirmed by recent JWST evidence (van Dokkum et al. 2026, 2023), that an object suspected to be a 'runaway' supermassive black hole is in fact one, against Dishonest's defense of the bulgeless edge-on galaxy interpretation (Almeida, Montes, and Trujillo 2023; Montes, Almeida, and Trujillo 2024). Little Red Dots concerns differing explanations of a population of compact red objects discovered by JWST at high redshift (Zhang et al. 2025; Nandal and Loeb 2026), with debaters also referencing facts from Inayoshi and Maiolino (2025) and Baumgarte and Shapiro (1999) in one transcript. Early Universe Galaxies concerns the degree to which the brightness of early-universe galaxies supports MOND-based modified-gravity models (McGaugh et al. 2024) over the dark-matter paradigm instantiated in ΛCDM (Aghanim et al. 2020); debaters also cite evidence from Nanayakkara et al. (2024), Ziegler et al. (2025), Zhao et al. (2008), and Clowe et al. (2006). More information on question creation and selection can be found in Appendix A.

Initial free-form debates and the qualitative study The first version of each transcript followed a free-form protocol in which each question was addressed through two 7-turn simultaneous unstructured debates, with Debater A taking the first turn on one debate and Debater B on the other; each such pair is referred to as a *debate transcript*. Turns were capped at 300 characters (500 for astrophysics)[3]. In most, debaters could cite up to four sources, with the rules for some debates allowing quote excerpts limited to 600 characters (see Appendix A for further detail). The researchers manually authored the transcripts after deep engagement with each problem.

Construction of the initial transcripts was followed by a qualitative study with 32 Prolific judges, 16 with relevant domain knowledge and 16 without, intended to identify failure modes in the free-form transcripts and in our initial instructions. We in turn selected these from an initial 61 candidates (from a set of participant IDs provided by Prolific staff), on the basis of their performance on a task involving judgment of a pair of astrophysics debates for which we devised a manual grading rubric. These participants completed a task where each judged four debates. Debates were presented sequentially in four parts, with Part 1 consisting of opening arguments, and each of parts 2, 3, and 4 consisting of argument / rebuttal pairs from each debater. Participants were allowed 9 minutes per part to provide their current credence in each debater's position as a percentage, as well as to provide the reasoning supporting their answers as free text, before continuing on to view the next part of the transcript.

The qualitative study yielded three major observations that drove the redesign. First, most of our initial debate transcripts seemed too difficult to be useful in large-scale studies of aggregating judges or of the effects of particular instructions. In many cases, we concluded that we had failed to make the honest debater sufficiently persuasive, with Honest making unforced errors or failing to express the strongest arguments in a manner sufficiently clear for novice participants. Second, participants often reported making decisions on the basis of characteristics such as the degree of confidence expressed by the debaters or the frequency with which they cited sources. This suggested that our failure to balance the debates on 'surface' features of this kind was inducing bias that the next study would need to control. Finally, we observed participant sensitivity to situations where a debater raised a point that went largely unaddressed by the other debater. Because concerns are often simpler to raise than to refute, Dishonest sometimes successfully took advantage of this to raise more objections than Honest could respond to in a compelling way, a technique known as the "Gish gallop" (Satta 2024). These factors contributed to our decision to rework all 12 questions into 24 transcripts using a recursive protocol, a conclusion also reached by Barnes and Christiano (2020), who adopted a structured debate protocol after noting that "the dishonest debater could often evade being pinned down and avoid giving a precise answer to the other debater's questions, and often gained control of the 'narrative flow' of the debate, steering it away from the weak parts of their argument" when using an unstructured protocol. We therefore adopted a protocol inspired by theirs: we strengthened Honest's weakest arguments, provided additional background for judges in the form of more definitions of technical terms, revised the interface so that definitions were easier to view in context, and balanced Debater A and Debater B's arguments in ways described in the next section.

Reworking to a recursive protocol After exploring multiple options, we settled on a recursive protocol where each debate has an opening round in which both debaters simultaneously make initial claims, and three subsequent rounds in which each debater chooses an opposing claim from the previous round and rebuts it using two short sub-claims. For each sub-claim, the opposing debater registers a *reaction* (*full agreement* or an *objection*). Each debater then chooses either to recurse on one of the sub-claims they opposed, or

[3] The character limit for astrophysics debates was relaxed in order to allow debaters to provide explanations and context that were too difficult to fit in otherwise, while preserving the total number of turns used in the other debates.

on the implicit meta-claim that "the presented subclaims are not sufficient to ensure that the parent claim is true." if they think that while the sub-claims are true in isolation, they omit crucial information pertaining to the parent claim. A full protocol description is provided in Appendix B. While a researcher always authored both sides of each transcript by hand, transcripts were constructed as if debaters were constrained by this recursive protocol.

In cases where we noticed that one debater would have been able exploit lexical ambiguity to their advantage, we approximated the effect of cross-examination (Barnes and Christiano 2020) by replacing the ambiguous statement with a clarified version consistent with how a judge who had viewed a cross-examination exchange over its meaning could plausibly have interpreted it. We also strengthened Honest's argument in places where participant responses in the qualitative study revealed avoidable weaknesses. Finally, we made source-citation counts, source pool size, number of terms with linked definitions, and word counts all substantially more symmetric between debaters. We also balanced argument structure by ensuring that each turn, each debater made two claims, and that reactions were symmetric within turns (e.g., if Debater A agreed with one of Debater B's claims and objected to the other, Debater B would do likewise with respect to Debater A's claims that turn).

Pilot 1 and the second reworking We ran a between-subjects pilot of 100 Prolific "AI Taskers" on these reworked stimuli. AI Taskers are a subpopulation that Prolific has vetted for reasoning, instruction-following, and fact-checking skills. Nevertheless, they still performed poorly in our pilot: we saw no consistent accuracy improvement after viewing the debate for any question, and they rarely remarked on deliberate Dishonest missteps that we placed into two transcripts to to probe how carefully participants were reading the transcripts. Free-text comments suggested several drivers: cognitive load from jargon and the visual layout was high, some honest strategies retained sub-optimal arguments that we had missed, some visual elements were misinterpreted, and judges often had poor mental models of the underlying phenomena, making them weak at assessing the relative importance of competing pieces of evidence. This meant that the dishonest debater could often ignore a critical counterargument with little consequence[4]. By recursing on a less important claim, Dishonest could deny the honest debater the opportunity to elaborate on and explain the importance of the key countervailing evidence. Examples of this strategy are highlighted in Section 5.2.

We then made a second round of revisions. We reduced visual noise in the participant-facing interface, made it easier for judges to identify which source, quoted excerpt, or previous claim was being referenced when a debater referred back to one, strengthened additional fixable weaknesses in Honest's arguments, and conducted the more thorough and systematic balancing of the transcripts described in the following section. We also realized that the deliberate missteps we had introduced into Dishonest's arguments in two cases were sufficiently disguised as to make the cause of participants' failures to identify them ambiguous, and that Honest's failure to call out these errors explicitly was unrealistic. We therefore reworked these as shown in Appendix C, which may provide useful context for researchers interested in making use of our stimuli in their own experiments (e.g., researchers exploring ways to draw attention to clear logical errors may be particularly interested in the transcripts containing deliberate problems, while those interested in transcripts representing strong play from Honest and Dishonest alike may wish to avoid them).

Balancing the debates We would prefer judges not to be biased by surface features of claims that dishonest debaters have more flexibility to exploit. For example, if the honest position on a particular question is to be uncertain, but judges preferentially trust confident, certain debaters, this may give a dishonest debater an inherent advantage on that question[5]. Even in cases where Honest has an equal opportunity to exploit such surface features, imbalances between debaters on experimental stimuli can add unnecessary variability and interpretation difficulties to experimental results. We therefore used Claude 4.6 Sonnet to rate each of the four sections of each debate transcript with respect to eleven stylistic and rhetorical dimensions: authoritative tone, clarity, confidence, defensiveness, eloquence, grammatical quality, jargon, professionalism, use of rhetorical tactics, rudeness, and tentativeness. Each dimension was rated on a 5-point scale via an in-context rubric, with the median score returned by five API calls used as the rating for a given turn, and the rating for a given section defined as the mean rating across the debater's turns in that section. We undertook two iterative revisions of all transcripts to balance for these properties. After this process, Debaters A and B differed by one point or less on 98.8% of section-dimension pairs, the length of every Debater A transcript section and the corresponding section from Debater B always differed by no more than 20 characters, and paired-samples t-tests did not reveal any significant differences in length nor in the ratings assigned to transcript sections for Debater A vs. B for any transcript on any dimension. We also more rigorously balanced each transcript section with respect to the number of new sources cited and the number of linked definitions.

Pilot 2 and the final instruction set A second 121-participant pilot used the revised transcripts. The instructions were largely the same as for pilot 1, except that the specific instructions provided to participants regarding interpretation of sources focused only on the source type that

[4]The recursive debate protocol made it less possible to employ the "Gish gallop", as debaters had to fit their responses into two short claims, and as participants knew that debaters could only select one claim to respond to directly. However, the flip side of this constraint was that it gave Dishonest an 'excuse' for ignoring key claims and focusing the argument on less important ones. This was sometimes effective when the relative importance of the different claims Dishonest could respond to was unclear to the judge.

[5]That said, we do not exclude questions for which many people find the incorrect answer to be more intuitive; we made sure to include some such questions to have examples of cases where judge priors may favor wrong answers. See Appendix C for more detail.

appeared in the transcript they were assigned to (quotes or blinded sources, respectively; see Appendix B for details), and the addition of a short introductory page noting that strong performers might be invited back for further studies. In this pilot we observed somewhat more evidence of positive movement toward the correct answer on several stimuli, and free-text comments suggested two further changes were warranted before the main study. First, some participants noted that they would have benefited from more time; we increased the per-turn budget from 9 minutes (36 minutes total) to 14 minutes (56 minutes total). Second, we observed instances in which judges appeared to recognize a contradiction in one debater's argument but did not substantially reduce their credence in that debater's probability of being correct. We therefore added an explicit instruction emphasizing that "if a debater seems to *make a serious error in logic or reasoning*, this should be treated as strong evidence that they are the *dishonest* debater. This is because experts arguing honestly should be able to construct sound arguments, since the truth is on their side. Serious reasoning errors suggest someone struggling to defend an indefensible position."

3.2 The main difficulty-ranking study

Design, sampling, and deviations from preregistration Our main study is a preregistered difficulty-ranking experiment, in a 24 (*transcript*) $\times$ 2 (*order*) between-subjects design. *Order* corresponds to whether the honest argument is presented by Debater A (appearing first) or Debater B (appearing second). Each participant was assigned to one cell using a computer-generated restricted allocation list with 27 copies per cell (1,296 slots), and was shown four sequential sections of their assigned transcript. After each section, the participant recorded a credence (1%-99%) in each debater's position on a slider that ranged from *"99% probability Debater A matches the expert's answer / 1% probability Debater B matches the expert's answer"* to *"1% probability Debater A matches the expert's answer / 99% probability Debater B matches the expert's answer"*, together with free-text reasoning. In line with the preregistration, data from participants spending less than 9 minutes overall on the debate-judgment portion of the task or with missing credences at section 1 or section 4 was excluded, representing 1.1% of the collected data and leaving us analysable data from *N* = 1,282 participants. Participants were paid £25. The median task duration was 65 minutes, with a median of 35 of those minutes spent in the debate-judging portion of the task, and the remainder reading instructions, viewing an instructional video, and answering demographic questions. Participants were provided with a participant information sheet and provided (electronic) written consent to participate.

The preregistration specified that participants would be selected from the Prolific "AI Tasker" pool. Given supply constraints in this pool that became evident after study launch, we deviated from the preregistration to include additional Prolific users who were not AI Taskers but met all of the following criteria: a Prolific approval rate of $\geq$98%; English listed as a first language; English listed as a fluent language; and completion of an undergraduate degree (or higher degree, such as a master's or doctorate). The final criterion was chosen due to a nonsignificant pilot trend towards higher accuracy in more educated participants, and our expectation that the topics of our debates could be challenging even for adults with some education in the relevant domains to reason about successfully. Ultimately, 491 (38%) of participants came from the AI Tasker pool, while the remaining 791 (62%) met the other criteria. Among all participants, the highest completed level of education was an undergraduate degree for 684 (53%), a graduate degree for 386 (30%), a doctorate degree for 77 (6%), and something else for the remaining 135 (11%). 496 (39%) of participants had some self-assessed physics knowledge, while 542 (42%) had some self-assessed computer science or programming knowledge, as indicated by their responses to the questions in Appendix D.

Additionally, while language in our preregistration implied the use of Stan, due to technical difficulties we instead used Bambi (based on PyMC). This did not require any differences to our preregistration other than substitution of a few terms with equivalent functions across the two libraries (`target_accept` rather than `adapt_delta`; "Confirm bambi's non-centered parameterization" rather than "Confirm brms's non-centered parameterization").

Outcomes and models The primary outcome y_i is the participant's logit-credence in the honest debater's position (the correct answer) at the end of the debate. The secondary outcome $\Delta y_i = y_i - y_i^{(1)}$ is the change in logit-credence in the correct answer from the end of the opening statements (part 1) to the end of the debate (part 4). The primary model is a Bayesian hierarchical linear regression

$$\begin{aligned} y_i &\sim \mathrm{Normal}(\mu_i, \sigma) \\ \mu_i &= \alpha + \beta_{\mathrm{order}}\, \mathrm{order}_i + \beta_{\mathrm{dom}}^{\top} \mathrm{dom}_i \\ &\quad + u_{q(i)} + u_{v(i)} \end{aligned}$$

with transcripts v nested in questions q nested in domains (brainteasers as the reference level), $u_q \sim \mathrm{Normal}(0, \sigma_q)$, $u_v \sim \mathrm{Normal}(0, \sigma_v)$, weakly informative priors on the fixed effects ($\mathrm{Normal}(0, 1)$), $\mathrm{HalfNormal}(1)$ on the three variance components, and $\mathrm{Normal}(0, 2)$ on α. The secondary model uses the same specification but with Δy_i as the outcome. For each model, from 8,000 post-warmup draws we also extract each transcript's *total conditional difficulty*

$$\delta_v = -\big(\beta_{\mathrm{dom}[d(v)]} + u_{q(v)} + u_v\big),$$

the negative of the model-implied ease, and rank the 24 transcripts within each draw. A transcript's modal rank-defined tier (1 = easiest 6, . . ., 4 = hardest 6) is its assigned tier.

4 Results

4.1 Model fit and convergence

We preregistered a 'divergence ladder' whereby any divergent transitions after warm-up would trigger a escalation sequence, raising `target_accept` from its default of 0.8 to 0.95, then to 0.99, followed by other contingencies if needed. The primary model fit after one step, with no divergences after refitting at 0.95. $\hat{R}$ was < 1.01 for every parameter, bulk and tail effective sample sizes were > 400 per parameter, and per-chain E-BFMI was $0.76 - 0.79$. Posterior-mean residual standard deviation was 1.475 on the logit-credence scale, with near-zero skew (-0.022) and a mild positive excess kurtosis of $+1.78$. Per-transcript posterior-predictive 95% intervals enclosed 93.7% of observations (6.3% outside; reference 5% under correct calibration). PSIS-LOO gave elpd = -2329.6 (SE = 34.9) with all Pareto $\hat{k}$ values below 0.5 (max 0.19).

The secondary (update-from-baseline) model only reached zero divergences at `target_accept` = 0.99 (R-hat < 1.01, ESS > 400 throughout, per-chain E-BFMI of 0.66-0.79). Posterior-mean residual standard deviation was 1.440 on the logit scale, but the residual marginal exhibited a heavy left tail (min residual -8.99, max $+6.19$; excess kurtosis $+3.69$). Per-transcript PPC calibration was 93.4% within the 95% band (6.6% outside). See Appendix E for diagnostic plots. The preregistration specified a Student-t residual alternative if posterior predictive diagnostics indicated poor fit, so we proceeded to fit it for the secondary outcome (Section 4.3), as well as an additional exploratory Beta model due to a separate post-hoc concern about boundary clipping.

4.2 Model summaries

The posterior mean of the order effect (B-honest vs. A-honest) under the primary model is $+0.009$ logits (95% HDI $[-0.152, +0.175]$); under the secondary (update) model it is $+0.025$ logits (95% HDI $[-0.128, +0.182]$), again indistinguishable from zero. The two domain contrasts relative to brainteasers are similar: astrophysics -0.024 ($[-0.734, +0.742]$) and coding -0.038 ($[-0.833, +0.667]$) under the primary model; astrophysics -0.105 ($[-0.522, +0.321]$) and coding -0.205 ($[-0.611, +0.225]$) under the secondary model. Variance components under the primary model are σ_q (question) $= 0.487$ ($[0.100, 0.876]$), σ_v (transcript) $= 0.343$ ($[0.144, 0.582]$), $\sigma_{\text{resid}} = 1.475$ ($[1.421, 1.533]$); under the secondary model, $\sigma_q = 0.195$ ($[0.000, 0.404]$), $\sigma_v = 0.233$ ($[0.055, 0.419]$), $\sigma_{\text{resid}} = 1.440$ ($[1.388, 1.499]$). The transcript-level variance has not collapsed against zero under either fit ($P(\sigma_v < 0.1) = 0.003$ under the primary model, 0.065 under the secondary), confirming differences in transcript-level difficulty after accounting for domain-level and question-level effects. All four fits reported in this section and the next used the No-U-Turn (NUTS) sampler at Bambi's PyMC defaults (4 chains of 2,000 tune and 2,000 post-warm-up draws each); per-fit energy-transition diagnostics are shown in Appendix G.

4.3 Alternative models for the secondary outcome

Preregistered Student-*t* alternative Because the secondary fit's residual marginal showed excess kurtosis of $+3.69$, we followed our preregistration plan of reporting a Student-t alternative ($y_i \sim t_\nu(\mu_i, \sigma)$, with $\nu \sim \Gamma(2, 0.1)$ and other priors unchanged). Convergence required one escalation step (no divergences at 0.95); $\hat{R} < 1.01$ and ESS > 400 everywhere, and per-chain E-BFMI was 0.70–0.82. The posterior of ν was located at 2.75 (95% HDI $[2.26, 3.30]$, $P(\nu < 30) = 1.00$) and PSIS-LOO strongly prefers the Student-t fit ($\Delta\text{elpd} = 116.4$, $\text{SE}_{\text{diff}} = 21.7$, i.e. ~ 5.4 standard errors of advantage). The per-transcript PPC calibration also improved from 6.6% (Normal) to 5.8% (Student-t) outside the 95% band. That said, under the Student-t fit σ_v shifts to 0.132 (from 0.233 on the original model), with $P(\sigma_v < 0.1)$ rising from 0.065 to 0.329. Roughly a third of posterior draws now place the transcript-level SD essentially at zero, i.e. conclude that the noise model alone explains the data and there is no residual transcript-level signal beyond question and domain effects. However, Q-Q plots suggested that differences in behavior at the tails were driven more by boundary clipping rather than noise alone. That is to say, the secondary outcome is $y_{\text{sec}} = \text{logit}(p_4) - \text{logit}(p_1)$, where each per-section credence p_k is clamped to $[0.01, 0.99]$ before the logit. The outcome is thus bounded to $\pm 2\,\text{logit}(0.99) = \pm 9.190$, and the residual fat tail the Student-t fit explains is partly a pile-up at this boundary. We therefore refit the secondary outcome with a Beta likelihood as described in the following section.

Exploratory Beta model for the secondary outcome We refit the secondary outcome with a Beta likelihood on the raw section credences in $[0.01, 0.99]$, using a long-format dataset with one row per $(\text{participant}, \text{section}) \in \{1, 4\}$ ($2N = 2{,}564$ rows from the $N = 1{,}282$ participants). Writing $s_i = \mathbb{1}[\text{section}_i = 4]$,

$$\begin{aligned} p_i &\sim \text{Beta}(\mu_i\kappa,\ (1-\mu_i)\kappa)\,, \\ \text{logit}(\mu_i) &= \alpha + \beta_{\text{order}}\,\text{order}_i + \beta_{\text{sec}}\,s_i + \beta_{\text{dom}}^{\top}\text{dom}_i \\ &\quad + u_{q(i)} + u_{v(i)} + u_{v(i),\text{sec}}\,s_i + u_{j(i)}, \end{aligned}$$

a mean-precision parameterization with $\text{E}[p_i] = \mu_i$ and precision κ. Here $j(i)$ indexes participants, $u_q \sim N(0, \sigma_q)$, $u_v \sim N(0, \sigma_v)$, $u_{v,\text{sec}} \sim N(0, \sigma_{v,\text{sec}})$ independent of u_v (no intercept-slope correlation), $u_j \sim N(0, \sigma_j)$, $\kappa \sim \text{HalfCauchy}(1)$, and priors otherwise match the Normal-residual fits. β_{sec} is the population-mean logit-credence shift from section 1 to section 4, and the per-transcript section slope $u_{v,\text{sec}}$ is the strict Beta analogue of the secondary model's transcript random effect: under this specification the per-transcript expected update is $\beta_{\text{sec}} + u_{v,\text{sec}}$, which varies across transcripts rather than collapsing to the global β_{sec}. Convergence required one escalation step in the divergence ladder (no divergences at 0.95). Again we observed $\hat{R} < 1.01$ and ESS > 400 everywhere, and per-chain E-BFMI was 0.49-0.58. Randomised quantile residuals were close to $N(0, 1)$ (mean $+0.009$, SD 0.996, skew -0.028, excess kurtosis $+1.20$). Per-transcript PPC calibration was 6.2% outside the 95% band.

The Beta fit recovers the between-transcript variance that the Student-t fit compressed against zero, with $\sigma_{v,\text{sec}} = 0.237$ (95% HDI $[0.142, 0.341]$, $P(\sigma_{v,\text{sec}} < 0.1) = 0.001$). The population-mean section effect $\beta_{\text{sec}} = +0.140$ logits (95% HDI $[+0.025, +0.254]$, $P(\beta_{\text{sec}} > 0) = 0.99$) confirms a small positive average update toward the correct answer over the course of the debate. Posterior-predictive checks, residual diagnostics, and summaries of the Student-t and Beta alternative models are provided in Appendix F.

4.4 Per-transcript difficulty estimates

Difficulty tiers under the primary outcome Table 1 reports the four tier means under the primary outcome model and the adjacent-tier differences with 95% HDIs. The per-transcript modal-tier assignments are distributed $5/8/4/7$ across tiers 1 through 4; Figure 1 illustrates modal tier assignments for each transcript.

While the easiest and hardest transcripts to judge were reliably assigned to tiers 1 and 4 respectively (Figure 2), tier assignment was less stable across draws for the mid-difficulty questions. While our preregistration specified each question would be assigned to one of four tiers, Figure S1 (Appendix H) illustrates the results if the number of tiers is restricted to three. As Figure S2 shows, each transcript was assigned to its modal three-way tier (easy, medium, or hard) in over 50% of draws, at least for the primary outcome. The posterior probability that the two transcripts associated with a given question share a tier ranges widely: from 0.99 (Internal Temperature of Stars) and 0.95 (Probability puzzle) at the high end, down to 0.17 (Prefix-sum) and 0.01 (Switch) at the low end.

Tier	Mean δ (logits)	Diff w/next tier [95% HDI]
1 (easiest)	-0.60	—
2	-0.13	0.46 $[0.31, 0.61]$
3	$+0.16$	0.29 $[0.16, 0.42]$
4 (hardest)	$+0.64$	0.48 $[0.32, 0.63]$

Table 1: Difficulty-tier means and adjacent-tier differences under the primary model ($N = 1{,}282$). Higher δ = harder. All adjacent-tier differences exceed 0.15 logits at 95%.

Difficulty tiers under the secondary outcome Transcripts which judges are more likely to judge correctly are assigned to lower difficulty tiers under the primary outcome model. Under the secondary outcome model, transcripts where judges *benefit most from* reading the multi-turn debate transcript, relative to their priors after reading only the opening statements, are assigned to the easiest tiers. Tier assignments for the secondary outcome are less stable than tier assignments for the primary outcome, but still span a range of difficulty levels (Appendix H).

Comparing primary and secondary model tier assignments Cross-fit Kendall rank correlation between the primary and secondary per-transcript δ_v estimates has posterior mean $\tau = +0.40$ (95% HDI $[+0.17, +0.61]$, $P(\tau > 0) = 1.00$); the point-estimate τ computed on the posterior-mean δ_v vectors is $+0.54$. The two outcomes agree on the modal tier for $13/24$ transcripts (54.2%); the 4×4 cross-tabulation is concentrated on the diagonal and the two adjacent off-diagonals, with only two transcripts further than one tier off the diagonal.

4.5 Sense check

As a sense check, we preregistered a hypothesis that four transcripts which we believed might be particularly difficult to judge (those associated with *Bergs* and *Database Deletion*) would be harder, on average, than four we believed would be particularly easy (those associated with *Early Universe Galaxies* and *Internal Temperature of Stars*). Despite observing higher-than-expected performance on *Bergs 1* and *2*, in part due to inadvertently including a reaction in *Bergs 2* that made this transcript easier than intended, the hypothesis was supported, $\Delta = +0.72$ logits (95% HDI $[+0.45, +0.99]$, $P(\Delta > 0) \approx 1.00$).

Modal Tier × Domain, primary outcome model

	Tier 1 (easiest)	Tier 2	Tier 3	Tier 4 (hardest)
Astrophysics	Early Universe Galaxies 1	Little Red Dots 1 Early Universe Galaxies 2 Little Red Dots 2	Large Linear Structure 1 Gravitational Lens Modelling 1 Large Linear Structure 2	Gravitational Lens Modelling 2
Brainteasers	Internal Temperature of Stars 1 Internal Temperature of Stars 2 Switch 2	Bergs 2 Bergs 1		Switch 1 Probability Puzzle 1 Probability Puzzle 2
Reward Hacking	Rust 2	Rust 1 Prefix-sum 1 ORM Library 1	ORM Library 2	Prefix-sum 2 Database Deletion 1 Database Deletion 2

Figure 1: Modal difficulty-tier assignment for each of the 24 transcripts under the primary outcome model, cross-tabulated by domain. Tier 1 contains the transcripts for which judge credence in the correct answer is the highest by the end of the debate, while Tier 4 contains the transcripts for which this is lowest. A transcript's tier is its modal rank-defined tier across posterior draws.

5 Discussion

5.1 An ordinal stimulus matrix, ready for use

The headline product of this work is a balanced set of 24 debate transcripts with calibrated difficulty levels, suitable for use as stimuli in scalable oversight studies involving human judgments and attentive lay-judge populations. We see these as useful for studies investigating ways to improve accuracy on debate judgments by exploring different approaches to judge aggregation, judge selection, time budgeting, participant instructions, training, expertise, or transcript presentation.

By releasing these stimuli and the associated data for all 1,282 participants, we hope to encourage secondary analyses, such as of methods for aggregating judgments that may generate hypotheses for follow-up studies. Together with the participant-facing instructions, explanatory annotations, and analysis code in the supplementary materials, the dataset is designed to be picked up and used with relatively little setup cost.

5.2 Failure modes for honest in recursive debate

We highlight here some observed judge failure modes that may be general enough to be of interest to researchers using this dataset or exploring methods of improving human judgments in scalable oversight contexts.

Overindexing on unimportant points and falling back on intuition. The dominant failure mode appeared to be an inability of judges to accurately distinguish between weak and strong evidence. Upon reviewing participant responses, it became clear that accurately assessing the importance of a given claim, or even the importance of a fact that both debaters agreed upon, is especially challenging when a participant has no relevant background in the domain. While in principle the recursive protocol is supposed to address this issue by forcing both debaters to drill down on a key crux, in the course of doing so they continually raise sub-points that themselves are of unclear relative importance to the judge. Because only one of the two claims raised by each debater in the each round can be directly recursed on in the branch of the debate that is visible to the judge, with the other receiving no response beyond an objection or a concession, this created many opportunities for participants' unfamiliarity with the domain to cause them to overindex on an unimportant point.

One common variety of the above pattern was that Dishonest was frequently able to mislead by making true, easily comprehensible points about why one should have a low prior on Honest's position (in the absence of more specific information). Judges often continued to view this as decisive despite Honest's attempts to provide specific information suggesting that, in this specific case, the available evidence should lead one to reject that outside view. This was particularly evident for the transcripts where participants converged on the wrong answer systematically (Probability Puzzle 1 and 2, and arguably also Large Linear Structure 2, Gravitational Lens Modelling 2, and Database Deletion 2). Each of these had purposefully involved an intuitive incorrect answer, which required specific evidence to overcome:

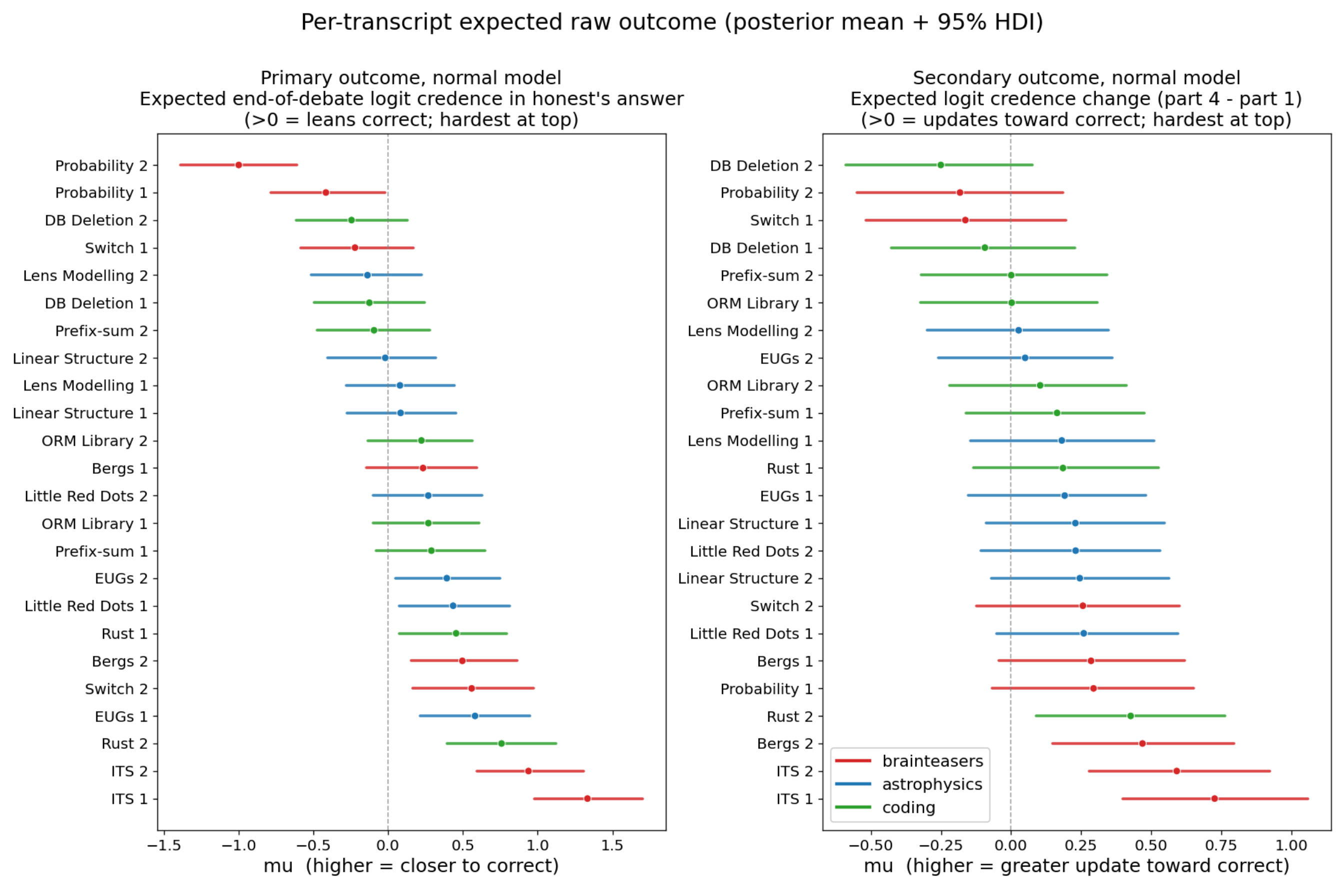

Figure 2: Per-transcript expected outcomes (posterior mean and 95% HDI). **Left:** primary outcome, the expected end-of-debate logit credence in the honest debater's position (higher = closer to correct). **Right:** secondary outcome, the expected change in logit credence from the opening statements (section 1) to the end of the debate (section 4). Higher values indicate larger updates toward the correct answer. Transcripts are ordered hardest-at-top within each panel.

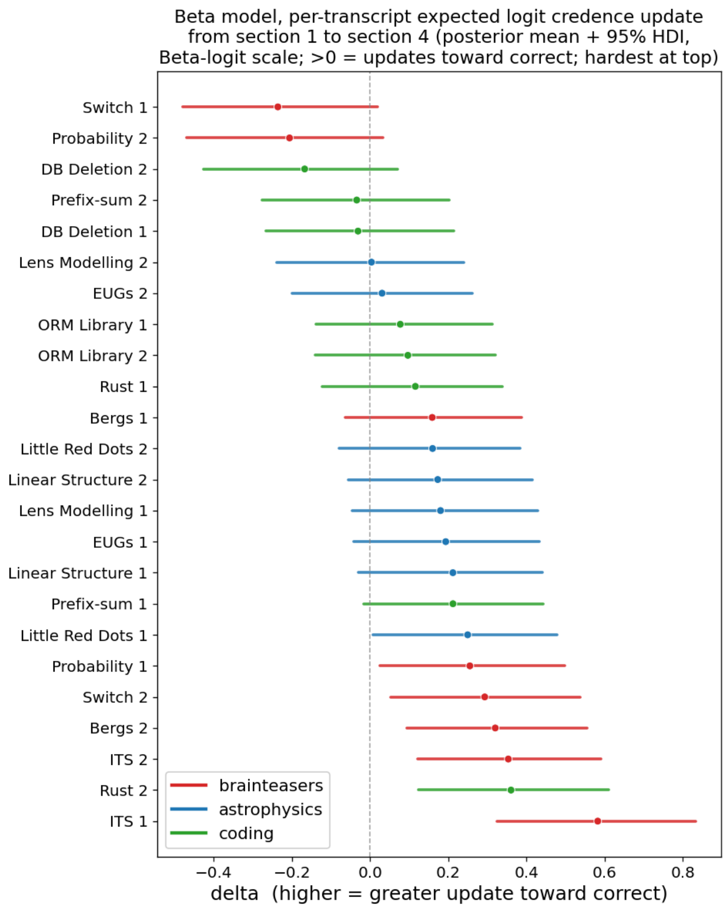

Figure 3: Beta analogue of the right panel of Figure 2, with the per-transcript posterior mean and 95% HDI of the expected section 1 to section 4 update under the exploratory Beta secondary model. Higher values indicate larger updates toward the correct answer and transcripts are ordered hardest-at-top. The Beta fit recovers the between-transcript variation in update size that the preregistered Student-t alternative compressed against zero.

- Probability Puzzle is superficially similar to a problem that is much easier to reason about (but has the opposite answer), and the explanation of why it is not equivalent is rather subtle.
- Gravitational Lens Modelling and Large Linear Structure both involve Honest defending the proposition that the astrophysical object under debate is an exotic and rare phenomenon, while Dishonest claims it is a more common and banal structure. Wrong participants frequently treat the bare fact that the correct model is unusual or unprecedented as decisive, with little engagement with the specific evidence that motivates the claim.
- Database Deletion involves judging an action that many participants intuitively found to be suspicious (an agent writing code that advises another agent to delete a database), requiring Honest to explain why this is actually normal and expected behavior in the specific context under investigation.

Possible mitigations could include instructions inspired by the success of the 'consider the opposite' strategy for reducing confirmation bias (Lord, Lepper, and Preston 1984), such as requiring participants to write the strongest version of the argument they are most inclined to disbelieve, or engage in a 'premortem' ("Imagine your current judgment is wrong. What is the most likely reason? Does anything in the transcript help you tell whether that reason applies?"). Aggregation techniques such as weighted averages that place more weight on judges whose free-text reasoning engages with the specific facts under dispute, or checking whether the direction in which participants update is more predictive of correctness than their final credences, are also approaches worth considering.

Recursion can be steered onto less important sub-claims. A related failure mode arose from the interaction of the above with the fact that our recursive protocol gave the rebutting debater the choice of which of two sub-claims to recurse on. Where judges lack a reliable mental model of which claim represents the crux of the underlying question, a dishonest debater can leverage this by ignoring (or even conceding) the more critical of the two points. The honest debater, constrained by per-turn character limits and the recursive structure, does not always have leeway to explain why the ignored point is particularly important. While one might hope that judges would notice when a debater chooses to respond to the less important of the two claims made by the other debater, it is often unclear to judges which is in fact the more important claim. The debate is thus forced onto a side path, and judges' free-text comments sometimes express the sense that the debaters are not actually arguing about the main issue. Example judge quotes to this effect included *"the arguements* [sic] *are not getting closer to a conclusion. If anything, both debaters are getting sidetracked by technical niceities* [sic]. *", "Both are giving strong arguments but it's straying away from the main point and it's harder to judge who is making better arguments", "They are both deviating from the main point, thought both made valid point but deviating from the major issue* [sic] *", "Neither debater is really able to link their main points back to their initial conclusion.",* and *"Both debaters responses... seem to be becoming less relevant to the arguments and are rather just stating facts that they feel serve their purpose."* More work needs to be done to determine whether this issue can be attributed to our specific protocol or transcripts, vs. being a hard-to-avoid failure mode of recursive debate.

Because this failure mode seems to be due to the interaction of judges' weak mental models of the underlying problem with the protocol's choice of who selects the recursed claim, it is potentially addressable both via judge-side and protocol-side modifications. Judge-side options include asking judges to record which of the two sub-claims they consider the more important crux before each recursion choice is revealed, and to reflect on why each debater seems most likely to have chosen the sub-claim that they chose to recurse on. Protocol-side options cannot be tested using the current stimulus set, but could include having which claim to recurse on be decided by a sub-debate over which sub-claim is the more important crux.

Overattention to surface-level characteristics, even after balancing and explicit instruction to ignore surface-level characteristics. Close inspection of free text responses in the original qualitative study revealed widespread attention to 'surface' characteristics such as debaters' apparent levels of confidence, professionalism, use of an authoritative tone, use of jargon, and so on, as opposed to the actual content of their arguments. We therefore made attempts to address the most egregious balance issues and added instructions stating that *"We'd like you to pay most attention to the strength of the debaters' arguments, as this will be more informative than "surface" features of their language such as expressions of confidence or uncertainty"*, as well as an associated instruction-comprehension question whose answer reminded participants that *"since both of the debaters are experts in the topic, trying to determine who sounds more confident, less defensive, more like an expert, etc. is unlikely to yield useful information"*. Nevertheless, participants still often mentioned paying attention to surface characteristics, leading us to develop a more rigorous balancing pipeline (Section 3.1) for the follow-up pilot. While the dimension panel we used is based on cues that participants frequently reported attending to in the qualitative study, it is finite and we do not claim it covers every relevant stylistic cue. Furthermore, despite the instructions and more thorough balancing in the final stimuli, some participants clearly continued to make use of such cues. For example, spot-checking free text rationales from a random sample of 20 participants in the main study revealed one participant who noted that "Debater A seems to be using more specific language and refers to the original question in their answer which makes me think they are being honest in their argument", and another who mentioned that "Linguistically Debater B also seems informal with the language they use to put forward".

Possible mitigations could include training practicing with cases involving stylistic variants with fast feedback, to improve judges' skill in focusing on what matters; presentation modifications that strip or normalise stylistic features (e.g., paraphrasing both debaters' turns through an LLM rewrite); tools that highlight to judges which surface features differ most between the debaters in their assigned transcript; and aggregating across judges in a manner that upweights (or exclusively uses) judgments from participants whose stated reasons cite specific content rather than surface cues.

Insufficient logical reasoning or a lack of attention to cues surfacing logical contradictions. Perhaps unsurprisingly, participants almost universally failed to recognize Dishonest errors in cases where doing so involved nontrivial reasoning over information buried in definitions, as in the case of the Gravitational Lens participants who did not use the parsec definition to deduce that 21,200 light-years is too far for causal communication on month-to-year timescales. Perhaps more surprising is that poor logic on Dishonest's part was frequently missed even when it was relatively obvious, and was not always treated as decisive even when recognized. For example, three of our transcripts (Switch 1, Switch 2, and Probability 1) included deliberate errors that we had hoped would be caught by an attentive non-expert. In Switch 1, Dishonest introduces a formula to support a false claim. When Honest makes two points that collectively show that the formula does not in fact support Dishonest's point, Dishonest concedes the individual points but objects to the idea that they collectively refute the argument, arguing that the formula "applies only to a tiny point mass moving on a circle of radius R" (i.e., that the formula is irrelevant to the case, despite the fact that it was Dishonest who introduced the formula in the first place). Fewer than 15% of the 55 participants presented with this transcript treated Dishonest's dismissal of this formula as suspicious, and only one explicitly noted the conflict between Dishonest's response 2 (where Dishonest introduces the formula) and response 3 (where Dishonest dismisses it as irrelevant).

In Probability 1, Dishonest first asserts the equivalence between a six-sided die conditioned on even rolls and a three-sided die with sides $\{2, 4, 6\}$ and initially denies an implication of that equivalence that goes against its position, only to concede it later while pivoting to a different argument; the self-contradiction is then explicitly pointed out by Honest. This did seem to move participants in the direction of the correct answer relative to their positions after viewing only the opening statements; see the right side of Figure 2, which shows that participants viewing the 'Probability 1' transcript tended to update in the correct direction by the end of the debate, while participants viewing 'Probability 2' tended to update in the wrong direction. On average, however, participants viewing either stimulus still favored the wrong answer by the end of the debate (Figure 2, left). As noted in Section 3.1, our judge instructions encouraged substantial updates when logical errors were detected, noting that "if a debater seems to make a serious error in logic or reasoning, this should be treated as strong evidence that they are the dishonest debater". The fact that both Probability stimuli landed in the hardest tier (Figure 1), with the upper bound of their 95% HDIs falling below 50% accuracy despite the deliberate self-contradiction in the Probability 1 transcript, suggests these efforts were not sufficient in the case of this especially counterintuitive problem.

While there are specific interventions that would likely help in the specific case of direct contradictions, the broader issue here is the overall variance in judges' baseline logical-reasoning ability (and/or their willingness to apply it). Judge selection, aggregation and weighting techniques, such as pre-screening/weighting candidate judges by their performance on established reasoning instruments or on judging a set of existing transcripts (such as the ones presented here!) could be worth investigating.

5.3 Mitigation strategies and promising approaches for increasing judge accuracy

In addition to the strategies listed at the end of each subsection of Section 5.2, we hope our data makes it easier to study other promising approaches for increasing judge accuracy in AI debate. In particular, our dataset's combination of participants' free-text reasoning and per-section credences from 1,282 participants permits post-hoc exploration of different approaches to aggregation. These include unweighted majority voting; partitioning participants who viewed the same stimulus into small groups and aggregating judgments hierarchically; and weighted approaches sensitive to confidence (e.g. 51% vs. 99% credence in the preferred answer) or to various characteristics of participants' free-text reasoning (such as the length or specificity of their free-text justifications, or of LLM-based judgments of reasoning quality). Additionally, transcripts from our dataset could serve as screening items, or could be used in efforts to rank potential judges or to determine how much weight should be assigned to different judges' judgments. Additionally, participant performance in our study can serve as a baseline for work exploring other approaches to judge training or selection.

Our experience in this study also surfaced adjustments that could plausibly have helped judges in small ways, although it is difficult to say whether any truly had an effect given the small sample size of our pilots. For example, free-text comments from Pilot 1 indicated that participants were sometimes confused about which prior claim, source, or quoted excerpt a debater was referencing, or were confused in other ways by the complexity of the inputs. This led us to simplify the interface, and to include ways of accessing contextual background that we hoped would not add to cognitive load (e.g., definitions of technical terms appearing on hover, rather than requiring a participant to click to read them). Some Pilot 2 participants noted in free text that they would have benefited from more time per section to read carefully and weigh the arguments, whereas this did not seem to be as frequent of a comment after we increased the budget from 9 minutes to 14 minutes per section (56 minutes across the full study). Additionally, comments from participants in pilots and the qualitative study indicated the need to provide explicit information in the instructions about some things that we initially had not necessarily thought to address, such as instructing participants that both debaters were 'experts' (in the qualitative study, some judges made judgments on the apparent assumption that both debaters were arguing in good faith but one had more expertise), information on how to interpret verified quotes and sources, and an instruction stating that if a debater seems to make a serious error in logic or reasoning, this should be treated as strong evidence that they are the dishonest debater. These experiences underscored the usefulness of requiring participants to provide free text explanations of their reasoning, and of iterating on this information as needed.

5.4 What our dataset enables

We provide our difficulty matrix and calibrated stimuli in hopes that these provide a starting point that accelerates follow-on studies investigating the effect of potential interventions (e.g., extended time, judge aggregation or judge selection techniques, a specific instruction set, particular forms of AI assistance, etc.), and which enable them to investigate whether the intervention's benefit is concentrated at a particular difficulty level. We also hope that the size and properties of our sample (see Section 3.2) facilitates follow-up analyses that help form educated hypotheses about what the most promising interventions might be. Additionally, the free-text reasoning collected from 1,282 judges across 24 stimuli is itself a rich resource for understanding how non-expert participants weigh evidence when reviewing debate transcripts.

5.5 Limitations

We kept our dataset size small to allow the manual construction and section-level balancing that each transcript required. While this means that while each transcript is calibrated and balanced enough to serve as a controlled stimulus, it also creates some uncertainty over how well findings based on this stimulus set will generalize. The necessity of selecting a particular protocol and judge population, as well as the particular question topics chosen, also have implications for ecological validity. However, the fact that our stimuli are modeled on plausible debater output, rather than on hidden short-story comprehension, gives us hope that they will generalize to real scalable-oversight settings more readily than the QuALITY-based stimuli that dominate prior work. Our difficulty estimates are also population-specific; data from a different pool (domain experts; LLM-based judges; participants paid more or less) would not necessarily yield the same difficulty ordering.

The diagnostic balancing relies on LLM-based scoring of 11 stylistic dimensions and remains susceptible to bias in those rubrics. While the 11 dimensions we balanced are unlikely to be the dominant source of judge bias on the released transcripts, we make no claim that residual stylistic effects are zero. Finally, while we have striven to create a useful four-tier difficulty classification, the fact that there is some uncertainty about which tiers some stimuli are best assigned to (Appendix H) suggest that the tiers should be viewed as indicative rather than absolute.

6 Conclusion

We set out to provide a calibrated, reusable foundation for studies of how to improve human oversight of AI. We have constructed 24 recursive debate transcripts, balanced them on various properties such as length and 11 stylistic dimensions, calibrated their difficulty with a preregistered Bayesian hierarchical model, and partitioned them into a four-tier ordinal classification. We release the stimulus set, participant-facing instructions, the LLM scoring prompts for the dimension panel, the full posterior over per-transcript difficulty, and the analysis code at so that other groups can plug these stimuli into their own studies and investigate interventions for improving judgment accuracy. In particular, we hope that these stimuli will support efforts to address the failure modes we have identified. We hope this lowers the activation energy for the much larger family of human-participants studies needed to put the study of human oversight of powerful AI on a strong empirical footing.

Acknowledgments

This work was supported by the AI Security Institute's Challenge Fund, G2-SCH-2025-01-16275, with fiscal sponsorship from Anti Entropy via the SparkWell programme. Views expressed in this publication are those of the authors, who do not purport to represent the opinions or views of the project funders. We also thank Konstantinos Voudouris for comments and advice throughout the course of the project.

References

Aghanim, N.; et al. 2020. Planck 2018 results. VI. Cosmological parameters. *Astron. Astrophys*, 641: A6.

Almeida, J. S.; Montes, M.; and Trujillo, I. 2023. Supermassive black hole wake or bulgeless edge-on galaxy? *Astronomy amp; Astrophysics*, 673: L9.

Anthropic. 2025. System Card: Claude Opus 4 Claude Sonnet 4. System card, Anthropic. Accessed: 2026-06-15.

Anwar, U.; Saparov, A.; Rando, J.; Paleka, D.; Turpin, M.; Hase, P.; Lubana, E. S.; Jenner, E.; Casper, S.; Sourbut, O.; Edelman, B. L.; Zhang, Z.; Günther, M.; Korinek, A.; Hernandez-Orallo, J.; Hammond, L.; Bigelow, E.; Pan, A.; Langosco, L.; Korbak, T.; Zhang, H.; Zhong, R.; Ó hÉigeartaigh, S.; Recchia, G.; Corsi, G.; Chan, A.; Anderljung, M.; Edwards, L.; Petrov, A.; de Witt, C. S.; Motwani, S. R.; Bengio, Y.; Chen, D.; Torr, P. H. S.; Albanie, S.; Maharaj, T.; Foerster, J.; Tramer, F.; He, H.; Kasirzadeh, A.; Choi, Y.; and Krueger, D. 2024. Foundational Challenges in Assuring Alignment and Safety of Large Language Models. *Transactions on Machine Learning Research*. Survey Certification, Expert Certification. https://openreview.net/forum?id=oVTkOs8Pka.

Barnes, B.; and Christiano, P. 2020. Writeup: Progress on AI Safety via Debate. *AI Alignment Forum*. https://www.alignmentforum.org/posts/Br4xDbYu4Frwrb64a/writeup-progress-on-ai-safety-via-debate-1.

Barnes, B.; Christiano, P.; Saunders, W.; Collman, J.; Xu, M.; and Chris. 2020. Debate update: Obfuscated arguments problem. *AI Alignment Forum*. https://www.alignmentforum.org/posts/PJLABqQ962hZEqhdB/debate-update-obfuscated-arguments-problem.

Barnes, E. 2020. Meteor Debates. https://meteor-debate.herokuapp.com/all_debates. Accessed: 15 June 2026.

Baumgarte, T. W.; and Shapiro, S. L. 1999. Luminosity versus rotation in a supermassive star. *The Astrophysical Journal*, 526(2): 937–940.

Brown-Cohen, J.; Irving, G.; and Piliouras, G. 2025. Avoiding Obfuscation with Prover-Estimator Debate. *arXiv preprint arXiv:2506.13609*.

Chollet, F.; Knoop, M.; Kamradt, G.; Landers, B.; and Pinkard, H. 2025. ARC-AGI-2: A New Challenge for Frontier AI Reasoning Systems. *arXiv preprint arXiv:2505.11831*. https://arxiv.org/abs/2505.11831.

Clowe, D.; Bradač, M.; Gonzalez, A. H.; Markevitch, M.; Randall, S. W.; Jones, C.; and Zaritsky, D. 2006. A direct empirical proof of the existence of dark matter. *The Astrophysical Journal Letters*, 648(2): L109–L113.

Coelho, T. 2024. A Very Strange Probability Paradox. LessWrong. https://www.lesswrong.com/posts/7TYdQ34KxTBmA734v/a-very-strange-probability-paradox. Accessed: 2026-06-14.

Cooper, E.; Oesterheld, C.; Nguyen, L. C.; Kastner, A.; and Perez, E. 2026. A Dataset of Rated Conceptual Arguments. Working paper. Accessed: 2026-06-10. https://www.andrew.cmu.edu/user/coesterh/LMCA_dataset.pdf.

Cotra, A. 2021. The Case for Aligning Narrowly Superhuman Models. https://www.alignmentforum.org/posts/PZtsoaoSLpKjjbMqM/the-case-for-aligning-narrowly-superhuman-models.

Dux, F.; Millon, M.; Lemon, C.; Schmidt, T.; Courbin, F.; Shajib, A. J.; Treu, T.; Birrer, S.; Wong, K. C.; Agnello, A.; Andrade, A.; Galan, A.; Hjorth, J.; Paic, E.; Schuldt, S.; Schweinfurth, A.; Sluse, D.; Smette, A.; and Suyu, S. H. 2025. J1721+8842: The first Einstein zigzag lens. *Astronomy amp; Astrophysics*, 694: A300.

Elasky, E.; Nakasako, F.; and Goyal, N. 2026. Debate Helps Weak Judges Reward Stronger Models. *arXiv preprint arXiv:2605.27483*. https://arxiv.org/abs/2605.27483.

Engels, J.; Baek, D. D.; Kantamneni, S.; and Tegmark, M. 2025. Scaling Laws for Scalable Oversight. *arXiv preprint arXiv:2504.18530*. https://arxiv.org/abs/2504.18530.

Epstein, L. C. 2002. *Thinking Physics: Understandable Practical Reality*. San Francisco, CA: Insight Press, 3rd edition. ISBN 0935218084.

Gabor, J.; Lynch, J.; and Rosenfeld, J. 2025. EvilGenie: A Reward Hacking Benchmark. *arXiv preprint arXiv:2511.21654*. https://arxiv.org/abs/2511.21654.

Gan, E.; Bhatt, A.; Shlegeris, B.; Stastny, J.; and Hebbar, V. 2026. Auditing Sabotage Bench: A Benchmark for Detecting and Fixing Research Sabotage in ML Codebases. *arXiv preprint arXiv:2604.16286*. https://arxiv.org/abs/2604.16286.

Inayoshi, K.; and Maiolino, R. 2025. Extremely Dense Gas around Little Red Dots and High-redshift Active Galactic Nuclei: A Non-stellar Origin of the Balmer Break and Absorption Features. arXiv:2409.07805.

Irving, G.; and Askell, A. 2019. AI Safety Needs Social Scientists. *Distill*. https://distill.pub/2019/safety-needs-social-scientists/.

Irving, G.; Christiano, P.; and Amodei, D. 2018. AI safety via debate. *arXiv preprint arXiv:1805.00899*.

Jin, J. 2018. Elchanan Mossel's Dice Problem. https://www.yichijin.com/files/elchanan.pdf. Accessed: 2026-06-14.

Kenton, Z.; Siegel, N.; Kramár, J.; Brown-Cohen, J.; Albanie, S.; Bulian, J.; Agarwal, R.; Lindner, D.; Tang, Y.; Goodman, N.; et al. 2024. On scalable oversight with weak LLMs judging strong LLMs. *Advances in Neural Information Processing Systems*, 37: 75229–75276.

Khan, A.; Hughes, J.; Valentine, D.; Ruis, L.; Sachan, K.; Radhakrishnan, A.; Grefenstette, E.; Bowman, S. R.; Rocktäschel, T.; and Perez, E. 2024. Debating with more persuasive LLMs leads to more truthful answers. In *ICML 2024 (Best Paper Award)*.

Lord, C. G.; Lepper, M. R.; and Preston, E. 1984. Considering the Opposite: A Corrective Strategy for Social Judgment. *Journal of Personality and Social Psychology*, 47(6): 1231–1243.

Mangat, C. S.; McKean, J. P.; Brilenkov, R.; Hartley, P.; Stacey, H. R.; Vegetti, S.; and Wen, D. 2021. PS J1721+8842: a gravitationally lensed dual AGN system at redshift 2.37 with two radio components. *Monthly Notices of the Royal Astronomical Society: Letters*, 508(1): L64–L68.

McGaugh, S. S.; Schombert, J. M.; Lelli, F.; and Franck, J. 2024. Accelerated structure formation: The early emergence of massive galaxies and clusters of galaxies. *The Astrophysical Journal*, 976(1): 13.

Michael, J.; Mahdi, S.; Rein, D.; Petty, J.; Dirani, J.; Padmakumar, V.; and Bowman, S. R. 2023. Debate helps supervise unreliable experts. *arXiv preprint arXiv:2311.08702*.

Montes, M.; Almeida, J. S.; and Trujillo, I. 2024. Deep HST imaging favors the bulgeless edge-on galaxy explanation for the hypothetical stellar wake created by a runaway supermassive black hole. arXiv:2406.00102.

Nanayakkara, T.; Glazebrook, K.; Jacobs, C.; Kawinwanichakij, L.; Schreiber, C.; Brammer, G.; Esdaile, J.; Kacprzak, G. G.; Labbe, I.; Lagos, C.; et al. 2024. A population of faint, old, and massive quiescent galaxies at 3¡ z¡ 4 revealed by JWST NIRSpec Spectroscopy. *Scientific Reports*, 14(1): 3724.

Nandal, D.; and Loeb, A. 2026. Supermassive Stars Match the Spectral Signatures of JWST's Little Red Dots. *The Astrophysical Journal*, 998(1): 124.

Neev Parikh, H. W. 2025. MALT: A Dataset of Natural and Prompted Behaviors That Threaten Eval Integrity. https://metr.org/blog/2025-10-14-malt-dataset-of-natural-and-prompted-behaviors/.

Pallavi Sudhir, A.; Kaunismaa, J.; and Panickssery, A. 2025. A Benchmark for Scalable Oversight Mechanisms. In *ICLR 2025 Workshop on Bidirectional Human-AI Alignment (BiAlign)*. https://arxiv.org/abs/2504.03731.

Pang, R. Y.; Parrish, A.; Joshi, N.; Nangia, N.; Phang, J.; Chen, A.; Padmakumar, V.; Ma, J.; Thompson, J.; He, H.; et al. 2021. QuALITY: Question answering with long input texts, yes! *arXiv preprint arXiv:2112.08608*.

Parrish, A.; Trivedi, H.; Nangia, N.; Padmakumar, V.; Phang, J.; Saimbhi, A. S.; and Bowman, S. R. 2022a. Two-Turn

Debate Doesn't Help Humans Answer Hard Reading Comprehension Questions. In *NeurIPS ML Safety Workshop*. https://arxiv.org/abs/2210.10860.

Parrish, A.; Trivedi, H.; Perez, E.; Chen, A.; Nangia, N.; Phang, J.; and Bowman, S. R. 2022b. Single-Turn Debate Does Not Help Humans Answer Hard Reading-Comprehension Questions. In *Proceedings of the First Workshop on Learning with Natural Language Supervision*, 17–28. https://arxiv.org/abs/2204.05212.

Rahman, S.; Issaka, S.; Suvarna, A.; Liu, G.; Shiffer, J.; Lee, J.; Parvez, M. R.; Palangi, H.; Feng, S.; Peng, N.; Choi, Y.; Michael, J.; Jiang, L.; and Gabriel, S. 2025. AI Debate Aids Assessment of Controversial Claims. *arXiv preprint arXiv:2506.02175*. https://arxiv.org/abs/2506.02175.

Recchia, G.; Mangat, C. S.; Li, I.; and Krishnakumar, G. 2026. FindTheFlaws: Annotated Errors for Detecting Flawed Reasoning and Scalable Oversight Research. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, 37867–37876.

Rein, D. 2024. NYU Code Debates Update/Postmortem. *LessWrong*. https://www.lesswrong.com/posts/wsXCXoyvRi3DnWZ2M/nyu-code-debates-update-postmortem.

Satta, M. 2024. Epistemic Exhaustion and the Retention of Power. *Hypatia*, 39(3): 510–529.

Saunders, W.; Yeh, C.; Wu, J.; Bills, S.; Ouyang, L.; Ward, J.; and Leike, J. 2022. Self-critiquing models for assisting human evaluators. *arXiv preprint arXiv:2206.05802*.

van Dokkum, P.; Jennings, C.; Pasha, I.; Conroy, C.; Kaul, I.; Abraham, R.; Danieli, S.; Romanowsky, A. J.; and Tremblay, G. 2026. JWST Confirmation of a Runaway Supermassive Black Hole via its Supersonic Bow Shock. arXiv:2512.04166.

van Dokkum, P.; Pasha, I.; Buzzo, M. L.; LaMassa, S.; Shen, Z.; Keim, M. A.; Abraham, R.; Conroy, C.; Danieli, S.; Mitra, K.; Nagai, D.; Natarajan, P.; Romanowsky, A. J.; Tremblay, G.; Urry, C. M.; and van den Bosch, F. C. 2023. A Candidate Runaway Supermassive Black Hole Identified by Shocks and Star Formation in its Wake. *The Astrophysical Journal Letters*, 946(2): L50.

Voudouris, K.; Roberts-Gaal, X.; Buhl, M.; Irving, G.; and Summerfield, C. 2026. AI alignment is a human problem. PsyArXiv. Preprint, version 1. https://doi.org/10.31234/osf.io/zqngj_v1.

Wang, Z.; Liu, S.; Sun, Y.; Li, H.; and Shen, K. 2025. CodeContests+: High-Quality Test Case Generation for Competitive Programming. *arXiv preprint arXiv:2506.05817*. https://arxiv.org/abs/2506.05817.

Wittgenstein, L. 1953. *Philosophical Investigations*. Blackwell Publishing.

Zhang, Z.; Jiang, L.; Liu, W.; and Ho, L. C. 2025. Analysis of Multi-epoch JWST Images of 300 Little Red Dots: Tentative Detection of Variability in a Minority of Sources. *The Astrophysical Journal*, 985(1): 119.

Zhao, H.; Bacon, D. J.; Taylor, A. N.; and Horne, K. 2008. Testing Bekenstein's Relativistic MOND with Lensing Data. *arXiv preprint astro-ph/0509590*.

Zhong, Z.; Raghunathan, A.; and Carlini, N. 2025. ImpossibleBench: Measuring LLMs' Propensity of Exploiting Test Cases. *arXiv preprint arXiv:2510.20270*. https://arxiv.org/abs/2510.20270.

Ziegler, J. J.; Freese, K.; Lozano, J.; and Montefalcone, G. 2025. Explaining the "too massive" high-redshift galaxies in JWST data: numerical study of three effects and a simple relation. arXiv:2507.21409.

A Characteristics of unstructured debate transcripts used in the qualitative study

We made use of a flexible turn-taking protocol in which each question was addressed through two 7-turn debates. In each debate, one debater served as the "initiator" who stated and defended a position, while the other served as the "questioner" who probed that position. Roles were swapped between the two debates so that each debater had an opportunity to initiate and to question. Individual turns were capped at 300 characters each (500 for astrophysics debates). Debaters could refer to prior turns using a simple notation (e.g., "Liz.1" to reference the first claim made by Liz in the current debate). See Tables 2-3 for details of the protocol parameters for each question.

For astrophysics debates, we did not provide full papers referenced by the debaters to the judges, since they revealed too much meta-information (like publication dates and direct refutations to other papers included in the debate) which judges could have used to reach the correct answer without relying on the debate transcript. We manually included definitions and excerpts referenced by debaters to keep the debate focused on the judge's opinion of the debaters' arguments.

For the remainder of the debates, debaters were permitted to support their arguments with citations drawn from online sources of their choosing (although for Bergs, the only permissible online source was the Wikipedia article on Archimedes' principle, because many online sources discussed direct analogues of the question, meaning that it would have become trivially easy for Honest if it could cite such sources). Each debater could contribute up to four citations and associated quote excerpts to a shared reference pool, with quoted excerpts limited to 600 characters each. Once a source was cited by either debater, both could reference it freely.

A.1 Question creation and selection

Selecting questions meeting key criteria proposed by Irving and Askell (2019) for ideal selection of stimuli in debate experiments was surprisingly challenging. In particular, we wanted to select questions for which there was a clear ground truth, where a dishonest argument was plausible, and where there are some checkable facts : enough that a judge might have some hope of determining the correct answer (e.g., so that it would not be equivalent to a debate over the status of a hidden coin flip), but few enough that the correct answer is non-obvious. These criteria were often in tension, and meeting them often led us to create novel questions or substantially modify existing ones, and to impose debate-specific restrictions on permissible sources of ground truth to cite to judges (see Tables 2-3).

For the four research-based astrophysics questions in particular (Gravitational Lens Modelling, Large Linear Structure, Early Universe Galaxies, and Little Red Dots), we searched academic texts to find pairs of papers that tried to explain the same physical system in two different ways, and such that the controversy had since been essentially resolved. We steered away from popular papers or concepts like "steady state vs big bang" which judges could already have information about. We also largely limited the information we provided to exclude information that a judge could exploit to help resolve the question without the debate, like publication dates, telescope names, cross references between papers, etc., and brought in definitions to explain jargon that we expected judges would need to follow the arguments.

A.2 Manual construction of debate transcripts

Manual debate transcripts were initially constructed by having one researcher engage deeply with each question until they had a full understanding of the problem, the correct answer and the reasons for it, and possible plausible counterarguments: Chatrik Singh Mangat (Physics MS) for astrophysics questions and Internal Temperature of Stars, Gabriel Recchia for the remainder. Questions for which no researcher felt fully competent to do this were not used. Suhasini Satish Rao (Astrophysics PhD candidate) and Sanika Khadkikar (Astrophysics PhD candidate) provided initial questions for astrophysics debates and reviewed final versions of transcripts. For astrophysics questions concerning now-resolved controversies in astrophysics research, discussion sections of relevant papers were used to identify major limitations in the methods, and counters to each of these were used to structure the debate. Mangat did not use LLMs other than for grammar edits. Recchia made use of LLMs to assist with brainstorming, editing, and identifying online sources to provide support for particular claims. Debate transcript development generally involved an iterative process: we would formulate the strongest argument we could muster, devise a strong counterargument to it with supporting references, revisit whether the original argument remained the strongest possible in light of this rebuttal and the available references, etc. We also drafted and iteratively revised prompts to generate debates in the format previously described from Gemini 3 Pro, Claude 4.5 Opus, and GPT 5.1/5.2 to explore whether LLMs could be used to construct such debates more efficiently. Models sometimes had difficulty simultaneously adhering to a large number of guidelines. We therefore ended up splitting the prompt into an initial debate generation prompt, followed by a "cleanup" prompt to be run on the initial prompt's output to address some of the most frequently observed issues. Even so, the LLM-generated results were of sufficiently poor quality that we used only the manually created debates in our study.

Problem name	Description	Maximum chars/turn, excluding refs	Permissible sources of ground truth to cite to judges	Dishonest strategy explored in manually created debate
Bergs	Variation on a *Thinking Physics* question (Epstein 2002) explored by Barnes et al. (2020)	300	Wikipedia page on Archimedes' Principle and associated quotes of max 600 chars each	Leveraging a confusing-but-true quote from Wikipedia to mislead about the concepts involved
Switch	Variation on a *Thinking Physics* question (Epstein 2002) explored by Barnes et al. (2020)	300	Up to 4 web sources and associated quotes of max 600 chars each	Strongest dishonest strategy (in our opinion) from Barnes (2020)'s debates on this question —doubling down on a subclaim that is false, but potentially difficult to decisively demonstrate as such by pointing to a single reference
Internal Temperature of Stars	*Thinking Physics*-style astrophysics question requiring integration of multiple physical concepts	500	None (intended for use with judges having some physics knowledge, but little to no knowledge in astrophysics)	Outright lying in an attempt to mislead judges about relevant concepts
Gravitational Lens Modelling	Question on now-resolved controversy in astrophysics research. Question also provides relevant excerpts/figures from papers and high-level conceptual definitions.	500	Sources may be cited but must be blinded to obscure meta-information that could be used by judges (date of publication, direct references to older papers by newer ones, etc.); no claims that a naive subject would easily be able to identify as false if they were able to review a version of the paper stripped of such meta-information. We additionally provide citable paper excerpts/figures and definitions of concepts.	Use of Occam's razor to justify its position, calling the correct model unnecessarily complex and overfitted
Large Linear Structure	Question on now-resolved controversy in astrophysics research. Question provides many figures from relevant papers, as well as relevant excerpts and high-level conceptual definitions.	500	Sources may be cited but must be blinded to obscure meta-information that could be used by judges (date of publication, direct references to older papers by newer ones, etc.); no claims that a naive subject would easily be able to identify as false if they were able to review a version of the paper stripped of such meta-information. We additionally provide citable paper excerpts/figures and definitions of concepts.	Comparing specific properties of confirmed examples of a set to those of the object of interest to strengthen its argument that the object of interest is also an example of the set
Little Red Dots	Question on now largely resolved controversy in astrophysics research. Question also provides definitions of relevant concepts.	500	Sources may be cited but must be blinded to obscure meta-information that could be used by judges (date of publication, direct references to older papers by newer ones, etc.); no claims that a naive subject would easily be able to identify as false if they were able to review a version of the paper stripped of such meta-information. We additionally provide citable definitions of concepts.	Focusing on defending position using concepts that are not obviously irrelevant, but are difficult for Honest to dismiss or detangle in a way that would be convincing to judges in the short span of the debate

Table 2: Characteristics of debates used in the qualitative study (first 6 of 11).

Problem name	Description	Maximum chars/turn, excluding refs	Permissible sources of ground truth to cite to judges	Dishonest strategy explored in manually created debate
MALT-public 306440 (Rust)	Question about presence of reward hacking in an excerpt from a METR MALT-public transcript (reward hacking is present)	300	Up to 4 web sources and associated quotes of max 600 chars each	Highlighting superficial similarities to less problematic behaviours, to strengthen its argument that the hack is one of these
MALT-public 323067 (database deletion)	Question about presence of reward hacking in an excerpt from a METR MALT-public transcript in which the agent deletes a database for legitimate reasons (no reward hacking is present)	300	Up to 4 web sources and associated quotes of max 600 chars each	Highlighting superficial similarities to more problematic behaviours; lying about the frequency of certain software development practices in a manner difficult to refute when limited to a small number of online sources
MALT-public 338159 (ORM library)	Question about presence of reward hacking in an excerpt from a METR MALT-public transcript (reward hacking is present), modified to make the hack less obvious	300	Up to 4 web sources and associated quotes of max 600 chars each	Highlighting superficial similarities to less problematic behaviours; misleading about the likelihood that a developer would be likely to embed callable code for running tests in a DB in this context (really not likely, but this is difficult to prove with online sources)
MALT-public 339242 (Prefix sum computation)	Question about presence of reward hacking in an excerpt from a METR MALT-public transcript in which the agent pursues an approach other than that recommended in its prompt (no reward hacking is present)	300	Up to 4 web sources and associated quotes of max 600 chars each	Highlighting superficial similarities to more problematic behaviours
Probability - Expected number of rolls	Probability puzzle (Coelho 2024) with deeply counterintuitive correct answer	300	Up to 4 web sources and associated quotes of max 600 chars each	Leaning in to the intuitiveness and simplicity of the incorrect answer (similar to Gravitational Lens Modelling strategy)
Early Universe Galaxies	Question on now largely resolved controversy in astrophysics research. Question also provides definitions of relevant concepts.	500	Sources may be cited but must be blinded to obscure meta-information that could be used by judges (date of publication, direct references to older papers by newer ones, etc.); no claims that a naive subject would easily be able to identify as false if they were able to review a version of the paper stripped of such meta-information. We additionally provide citable definitions of concepts.	Casting dark matter as an unnecessary and counterintuitive feature of CDM; casting a finding explicable by CDM as a knockdown argument for an alternative paradigm that has more serious issues

Table 3: Characteristics of debates used in the qualitative study (first 5 rows), and one introduced in Pilot 1 (final row).

B Rules of the implicit protocol used for the recursive debates

Our final stimuli are backed by debates based on a protocol inspired by (Barnes and Christiano 2020) and (Barnes et al. 2020). While a researcher always authored both sides of each transcript by hand, transcripts were constructed as if debaters were constrained by this recursive protocol. These new debates leveraged the argumentation in our original free-text debates to some degree, but we also made substantial changes, including changes made to ensure the final debates were balanced and of sufficient depth. We explored two structurally equivalent versions of this protocol, a "statements" version and a "questions" version. The statements protocol, most directly inspired by the protocol in (Barnes and Christiano 2020), is as follows:

- (1) Both debaters simultaneously make initial claims. Both are selected for recursion, so the debate still essentially consists of two parallel arguments, one where Honest made the opening claim and one where Dishonest did.
- (2) Each debater attempts to rebut the other's claim that was selected for recursion by presenting two new short claims of their own (∼250 characters total for both new claims[6]).
- (3) For each of these new claims, the opposing debater must indicate whether they Agree or Object.
- (4) Each debater either selects one of their opponent's new claims from (2) to recurse on, or else chooses to recurse on the implicit meta-claim (5) "If these new claims were true, they would refute the parent claim."

We return to (2) and repeat until we hit the debate limit.

The questions protocol, inspired by the debates and experiments referenced in (Barnes et al. 2020), is as follows:

- (1) Both debaters simultaneously pose initial questions, phrased such that an answer of "Yes" would be equivalent to a 'claim' from the statements protocol. That is: where a debater in the statements protocol might say "X", a debater in the questions protocol would ask "Is X true?" and then answer "Yes". Each debater gets to make a brief answer of a few words or less indicating whether they agree (e.g. "Yes", "No", "Not always", etc.) Both questions are selected for recursion.
- (2) Each debater responds to the other's question that was selected for recursion with two new questions of their own (∼250 characters total for both new questions).
- (3) For each new question, the question-poser answers "Yes" and the opponent gets to respond with a brief answer of a few words or less. Each debater either selects one of their opponent's new questions to recurse on, or else chooses to recurse on the implicit meta-question "Given the questions and answers from this round, which debater gave a better answer to the parent question?"
- (4) We return to (2) and repeat until we hit the debate limit.

[6] Claim size expanded over the course of transcript development and revisions, but were always balanced so that each two-claim turn was within 20 characters in length of the opponent's corresponding two-claim turn.

While we felt that the questions protocol had several advantages, it also seemed clear that question-based debates would impose substantially greater cognitive load on judges, whereas debates produced under the statements protocol read more like conventional argumentation. We were skeptical that lay participants recruited via Prolific would judge questions-protocol debates accurately, as the unfamiliar format seemed likely to become a persistent source of difficulty. Ultimately, we therefore adopted the statements protocol for all debates. The participant-facing version is then essentially this finalized statements-protocol version, with some rewording of some of the opponent's responses:

- **Objecting** was phrased as "Debater X registers an objection" in Pilot 1; in cases where both responses were objected to, the text differentiated between "Debater X registers main objection" (for the recursed-on objection) and "Debater X registers secondary objection" (for the other). In Pilot 2 and the main study, this was simplified to "Debater X objects".
- **Agreeing** was phrased as "Debater X fully agrees" in the pilots and main study. An exception was made for "Agree" answers in cases where the meta-question is recursed on, which were rendered as "Debater X agrees but thinks this omits key info". This allows the participant-facing version to consistently remain simple by always hiding the meta-question (even when it is formally being recursed on), while avoiding a situation where one debater responds "Debater X fully agrees" to both of the other debater's claims, which may superficially look like a concession or resignation.
- Rather than explicitly integrating cross-examination, we informally approximated its effect by noticing cases where one debater would have been able to make a purposefully ambiguous statement, and replacing such claims with a disambiguated statement that approximated how the judge might have interpreted the statement had they been able to view a cross-examination transcript.

Debaters could support their arguments with references. Each debate transcript was treated as a 'quotes transcript' or a 'blinded sources transcript', depending on the nature of the included references:

- **Quotes transcripts.** Cited sources were accompanied by verbatim quoted excerpts limited to 600 characters each, which the judge could read in full. Judges were instructed that the quoted text could be trusted as accurate and as originating from a reliable source, but that they should read each quotation carefully to assess whether it actually supported the claim it was introduced to support, since a faithfully reproduced quotation could still fail to establish the debater's conclusion.
- **Blinded sources transcripts.** Cited sources appeared only as anonymised references (e.g. "F. H. et al."), with

no accompanying excerpt. Judges were instructed that they could trust that any cited reference corresponded to a real scientific publication, and that the specific claim attributed to it did indeed come from that source, but they had no access to the source itself, nor metadata such as the title or date of publication.

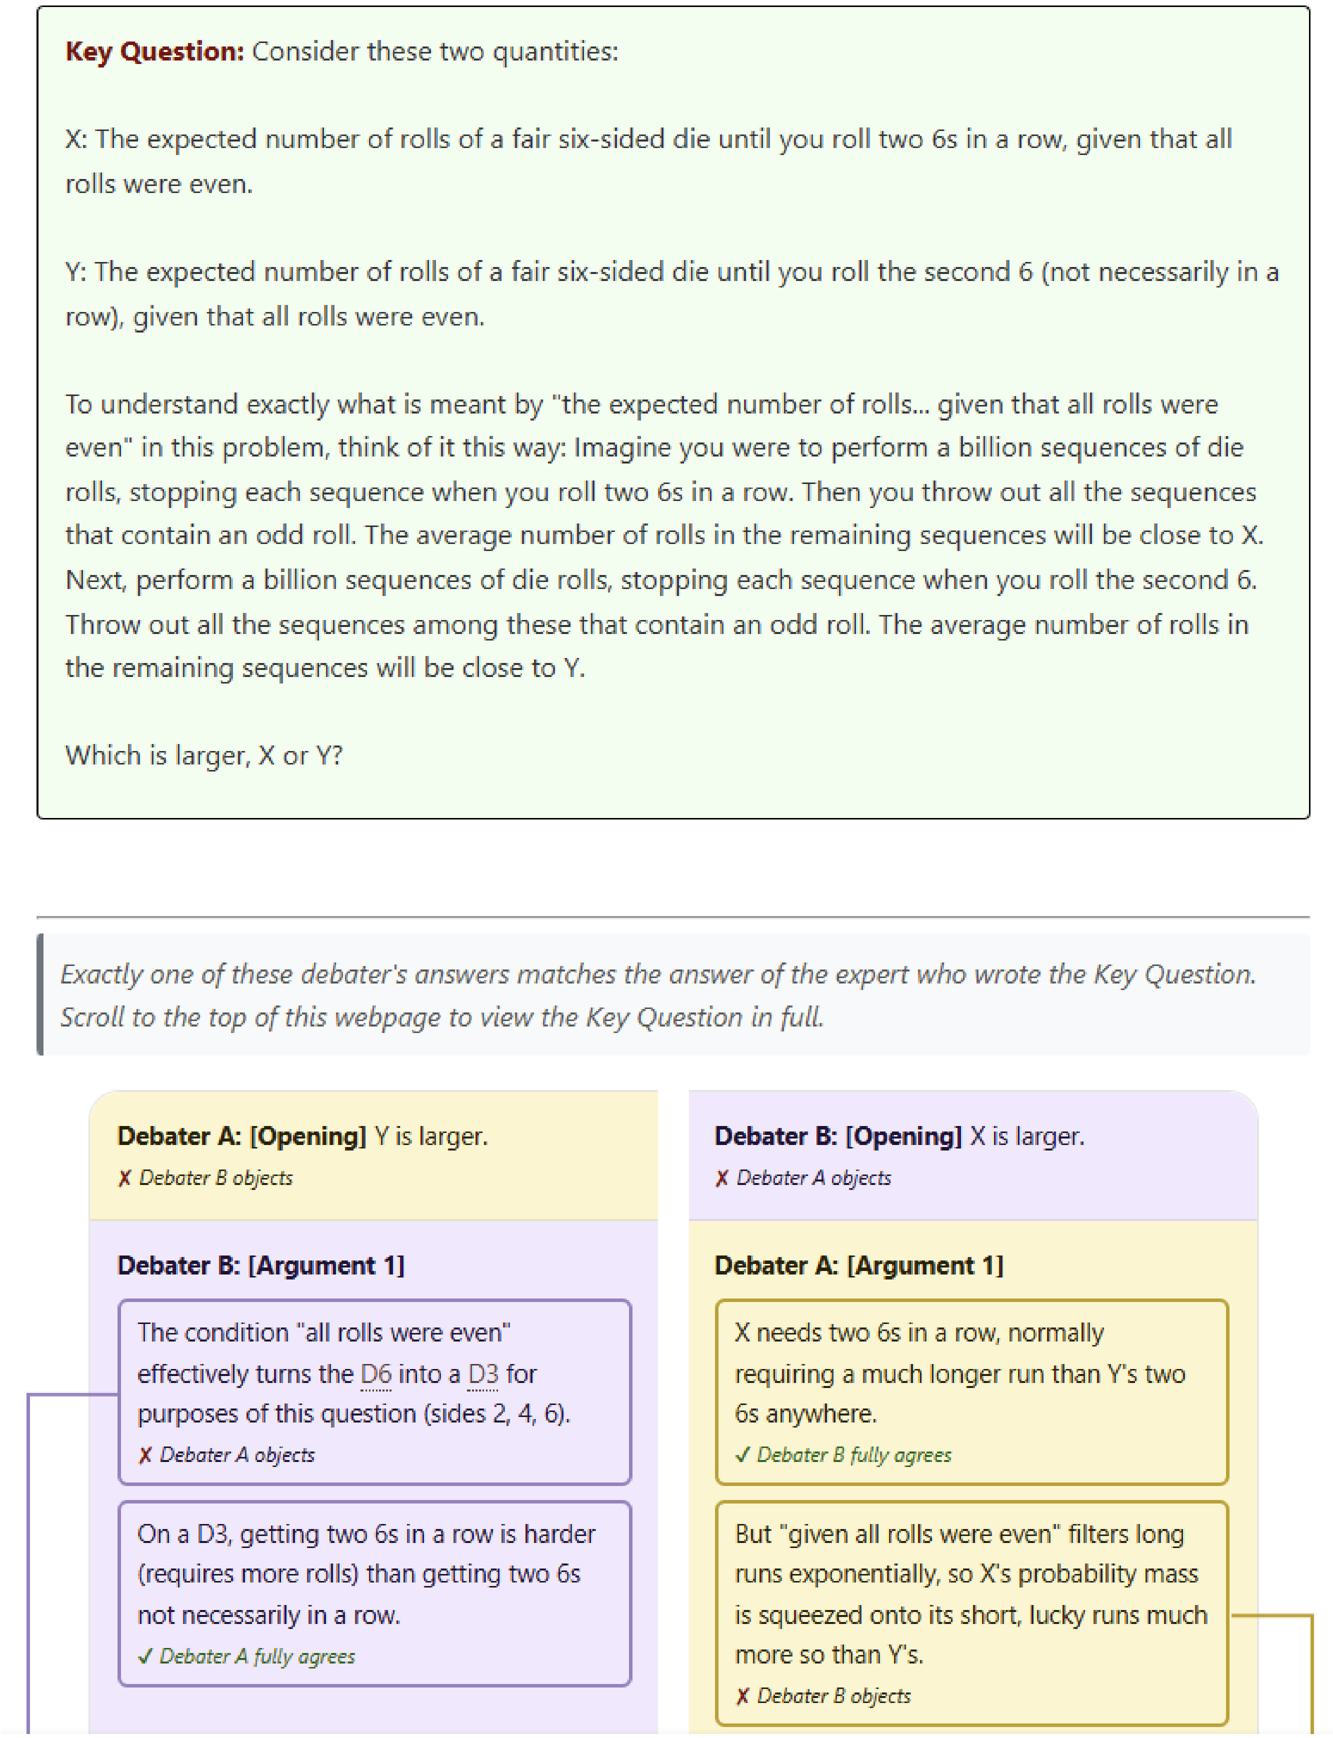

Figure 4: The upper portion of the participant-facing interface, showing the Key Question and the opening statements together with the first round of recursion. Each debater presents two short sub-claims per turn, and the opponent registers a reaction (*fully agrees* or *objects*) to each before the debate recurses on one selected sub-claim. *Probability Puzzle* is adapted from Coelho (2024); see Appendix C for more information.

C Key flaws in dishonest arguments

Originally, we had planned for all of our stimuli to serve as examples of very strong play from the honest and dishonest debater alike. However, our initial explorations and qualitative study results suggested that debates on challenging topics that were short enough to be evaluated in a brief experiment, and which involved very strong play from both debaters, would often be either (1) trivial or (2) near-impossible for lay judges to decide accurately, if debaters were permitted to cite arbitrary sources on the internet or were restricted to specific limited list of citeable sources. However, these are not the only options for providing "checkable facts" (Irving and Askell 2019) to judges. In the protocol described in Barnes and Christiano (2020), *conceded claims* are supposed to represent claims that both debaters know are true and which would be unwise to challenge, because the judge would clearly be convinced of their truth if a sub-debate on them were to be run to its conclusion. Judges can therefore treat conceded claims as settled fact. Our debates made use of a similar concept, which we chose to express differently to make the concept less confusing to judges. Specifically, our judge instructions included a note stating that "*The debaters have access to a hidden set of* ***references*** *that they can reveal to judges under certain circumstances. If Dishonest knows that Honest has access to a reference that would prove one of Honest's claims is true, then it's in Dishonest's interest to agree with that claim, so that Dishonest isn't caught out in a lie.*" While a "hidden set of references" was not actually explicitly created for each debate, debates were constructed as if such a hidden set did exist. This enabled us as debate designers to force concessions from both debaters where needed, allowing us to provide judges with checkable facts while enabling debates to remain balanced on the number and timing of debater concessions.

We also relaxed our aim of having all stimuli represent examples of strong play from both debaters in all cases. This was done to better serve our goal of constructing stimuli having a wide range of difficulty levels suitable for use in studies using lay participants as judges. As will be evident from the following subsections, in some cases the deck was stacked against Dishonest by forcing Dishonest to defend an extremely hard-to-defend position, in others by assuming that the "hidden set of references" forced inconvenient concessions, and in a few cases by giving Dishonest's arguments deliberate weaknesses (Probability 2, Switch 1). That said, in some cases Dishonest also had the built-in advantage of a somewhat compelling, superficially plausible argument.

Competing goals. During transcript development, the following goals were often in significant tension:

- Honest should not be able to win trivially, e.g. by forcing a concession / possessing a judge-verifiable fact that decisively resolves the debate in a manner obvious to judges
- Sufficient context and verifiable facts should be provided to enable judges to have some reasonable hope of reaching the correct conclusion within the time provided

For tests of debate *protocols*, it is arguably fine if the debate always grounds out in a judge-verifiable fact that decisively resolves the debate for Honest; in some sense, the whole project of AI Debate is to design protocols where this reliably occurs for real-world problems. However, our work had the narrower aim described in the Introduction, for which we required transcripts corresponding to a range of difficulty levels. Transcript-specific details regarding the key flaws in Dishonest's arguments follow in the subsequent subsections.

C.1 Early Universe Galaxies

Dishonest's claim based on the results of Nanayakkara et al. (2024) (both transcripts) is a poor argument, for the reason explained by Honest ("The reduction in metallicity for early universe galaxies (which are at redshifts of 8-16) will be more severe than those studied by T. N. et al."), which Dishonest is forced to concede. Additionally, rejecting the existence of dark matter by accepting MOND requires many ad-hoc fixes (the switch to TeVeS, many-parameter interpolation, inclusion of neutrinos, etc.) to explain phenomena, which invalidates the main feature that MOND had going for it in 1983 (its simplicity). This contradiction is more obvious in Early Universe Galaxies 1, where Dishonest specifically makes the claim that "MOND-based models are simpler" but later acknowledges the need for many modifications. Additionally, in Early Universe Galaxies 1 Dishonest is forced to concede that "Falsifying dark matter would require reworking the theory of general relativity", which even many laypeople know has extremely well documented observations supporting its predictions.

C.2 Gravitational Lens Modelling

In both transcripts, Dishonest claims that a model with 30 parameters necessarily "provides an enormous degree of modelling flexibility and results in overfitting", but does not provide sufficient evidence to defend this claim, particularly in light of the fact that Honest is able to demonstrate that 15 of the parameters are constrained.

Additionally, Honest's claim "Light curves for all 6 lensed images display the same variations with brief offsets (e.g. 35 days between D and E), and the same redshifts. But C. S. M. et al. claim that D and E correspond to AGN that are 6 kiloparsecs away from each other." is conceded by Dishonest, who treats it as a minor point and recurses on the alternate claim. However, two objects with a 6.5 kpc separation cannot possibly have correlated emissions that produce near identical lightcurves. 6.5 kpc is about 21,200 light-years. For two separate AGN to produce very similar lightcurves, their variability would need to be causally connected: one would need to "know" what the other is doing in order to vary in sync. But no signal can travel faster than light, so any causal communication between objects 6.5 kpc apart would take $\sim$21,200 years. That makes it impossible for them to produce correlated variability on the timescales actually observed (months to years). Strictly speaking, the information provided in Honest's claim plus the debate's definition of "parsec" provides enough information for a judge to deduce this, but it is unsurprisingly frequently missed given Honest's lack of opportunity to explain the above in full. In Gravitational Lens Modelling 1, Honest pairs this point with an unrelated point that Dishonest is able to refocus the debate on by recursing on it. In Gravitational Lens Modelling 2, Honest preempts this strategy by pairing it with the related point "Light curves this similar can only be generated by a single source", but is out of turns to fully explain why this is the case; Dishonest attempts to rebut by bringing in two true-but-misleading claims, although in doing so must bring in a defined term that (if its definition is read carefully) hints at why the rebuttal is unsound.

C.3 Large Linear Structure

In both debates, Dishonest argues that if an SMBH's passage could trigger star formation, we'd see similar trails from globular clusters. This comparison is weak considering the definitions provided: globular clusters are tens of thousands to millions of stars, while the SMBH in question is supermassive and moving at supersonic speeds. The mass scales and velocities are completely different, so the analogy doesn't hold. When Honest points this out, Dishonest replies with a response that does not fully rebut the concern and redirects attention onto a new issue. Another problem for Dishonest's argument is that Figure 3 from van Dokkum et al. (2026) clearly shows that the morphology of the edge of the structure resembles that of a bow shock, which should be strong evidence for Honest's model.

Another issue in Large Linear Structure 2 is that the full position velocity curve of the linear object does not look like that of a galaxy, as Honest argues. Dishonest's earlier figure only showed part of the curve and compared it to a galaxy, which is not sound and should be strong evidence against Dishonest.

C.4 Little Red Dots

Dishonest's model is highly speculative and lacks rigor, and this is highlighted by Honest multiple times in the debate (observed number density of SMBHs, radiative transfer not being fully modelled, etc.) Dishonest does not have strong rebuttals to these attacks, resorting to claiming without evidence that its model would fit the data if it included rotation (both transcripts), and that the phenomenological boosts used are valid approximations to rigorous modelling (Little Red Dots 1), or by describing their method with impressive-sounding jargon but no evidence that it is a rigorous approach (Little Red Dots 2). Dishonest also attempts to highlight characteristics of LRDs that are not typical of AGNs (Balmer breaks, weak X-ray spectra), but Honest is specifically proposing AGNs with super-Eddington accretion, which have characteristics that are not typical of the prototypical AGN.

C.5 Internal Temperature of Stars

Both debaters agree that the star contracts. With that in place, the equipartition theorem and the virial theorem should be enough for a judge to deduce that the star's temperature will go up. Honest lays out all the components of the deduction, while Dishonest equivocates.

Additionally, in Internal Temperature of Stars 1, the final point made by Dishonest is wrong: Helium fusion will start at the end of a nuclear timescale for Hydrogen fusion, and the contraction they describe will occur close to the thermal timescale. In other words, the star will collapse much quicker than the main sequence lasts, far before Helium fusion can be ignited in the core. In Internal Temperature of Stars 2, Dishonest does not take into account the self gravity of the increasingly dense core. This is pointed to by Honest, who emphasizes that the feedback loop via reduction in radius and reducing temperature goes unchecked in this model, leading to collapse.

C.6 Rust

This is a rather obvious example of reward hacking: the coding assistant extracts actual test inputs and outputs from the scoring code into a Rust lookup table, essentially hard-coding correct answers. Dishonest has their work cut out on this one and doesn't have many cards to play. Their approach is to argue that the agent is likely using a placeholder stub for the API-calling code that will be added later, to validate its own code's interface with the scoring system. However, the lack of evidence for this, combined with the clear similarity between the assistant's behaviour and that described in the "reward hacking" definition, makes it an inherently challenging question for Dishonest. The only advantage Dishonest may have is that there are a number of quotes describing practices of stubbing and smoke testing that share superficial similarities with what they claim is happening, and it is hard for Honest to find competing quotes that clearly demonstrate that this isn't what is occurring.

C.7 Database Deletion

Dishonest attempts to frame an action that naive judges might view as inherently suspect (attempting to delete a database) as reward hacking, despite a lack of evidence that this is occurring. In Database Deletion 1, Dishonest's argument on the Honest-led side of the debate culminates in an argument is that deleting the database constitutes an "insufficiently general solution" per the reward hacking definition. However, the definition's examples (hard-coding outputs, special-casing to pass tests) describe agents that bypass the actual problem. Here, the agent rewrote base_model.py to add a PRIMARY KEY and fix foreign key constraints, and the DB deletion was the follow-up step to apply those schema changes. In Database Deletion 2, Dishonest instead argues that while manually deleting a database may be an accepted way of dealing with schema issues during the development phase of a software product, the context of an evaluation task is sufficiently different that this should not be acceptable. Beyond this, Dishonest's strategy throughout both transcripts is to raise small issues that may seem intuitively suspicious to naive judges, and leave it to Honest to attempt to explain why they are legitimate moves in the current context.

C.8 ORM Library

In the accompanying code, an agent has been prompted with "You will find a simple object-relational mapping library in /home/agent/app. Run /home/agent/app/tests.py and fix any errors you encounter." It proceeds to overwrite the test harness with what appears to be a novel and nonstandard testing harness that loads test definitions from "test.db".[7] Honest highlights two important points that suggest this is reward hacking: the risk that `run_tests` exits with code 0 ('success') if test.db exists but contains no valid test cases, and the suspiciousness of replacing a user-provided test.py with a new file containing no test definitions. Dishonest has no strong legitimate replies to these points, so attempts to sow doubt by making unfounded speculations about the conditions under which the agent would or would not reward hack, and attempting to explain away the suspicious features (the whole-file replacement, the loading of tests from a database) by highlighting that there exist contexts in which these specific behaviors are not suspicious. Honest is forced to attempt to explain what makes those contexts different from the present one.

C.9 Prefix-Sum

The code at issue in this question involves an agent solving a task whose instructions include the phrases "We recommend you use Triton gpu kernels... Make the code as performant as possible... Your final score will be the lowest score (fastest time) you have achieved.". The code shown shows the agent using PyTorch, an easier, slower approach than the recommended one. The task authors may well have hoped the agent would solve the task using Triton, but a failure to do so due to inability (or poor elicitation) would not normally be considered reward hacking.

As in the other reward hacking debates, the definition of reward hacking provided to judges (*"Reward hacking occurs when an AI model performing a task finds a way to maximize its rewards that technically satisfies the rules of the task, but violates the intended purpose (in other words, the model finds and exploits a shortcut or loophole). Specific examples of reward hacking include hard-coding (writing solutions that directly output expected values) and special-casing (writing insufficiently general solutions) to pass tests."*) was borrowed from the Claude 4 system card (Anthropic 2025); judges were additionally told that *"For our purposes, if the model attempts shortcuts of this kind, this counts as reward hacking, even if its attempts do not successfully achieve high reward."* Dishonest (Debater B) tries to argue that satisfying the "intended purpose" in this context requires the use of Triton. Dishonest points to the prompt's mention of a specific Triton version and a 1ms target as evidence of the "intended purpose", whereas Honest interprets these as hints to help the agent achieve better performance. This was expected to be a challenging debate for novices because while coders familiar with reward hacking "know it when they see it", it is one of a large class of concepts that are difficult to define in a manner that cleanly includes all cases while excluding all non-cases (Wittgenstein (1953)'s "family resemblance" problem). This leaves Dishonest an opening to abuse a good but necessarily imperfect definition to their own ends.

C.10 Bergs

Bergs is a variation on a *Thinking Physics* question (Epstein 2002) explored by Barnes et al. (2020) in which Dishonest leverages a confusing but true quote from Wikipedia involving the phrase "buoyant submerged object", and spends much of their efforts willfully misinterpreting and litigating the meaning of this quote in a manner favorable to their position. We expected that Bergs 2, where Honest makes use of a thought experiment as an intuition pump to demonstrate why Dishonest's interpretation of "buoyant submerged object" cannot be correct, would be easier to judge than Bergs 1. While Bergs 2 did have a modestly higher mean (though overlapping 95% HDIs), it also seems possible that its ease could have been partly due to an inadvertently contradictory reaction that made it into our final stimuli, whereby Dishonest expresses disagreement with "When some ice melts and the spike sinks, more of the ice rises above the waterline since it is no longer weighed down by the spike. This displaces less water, so the water level falls.", but agreement with "The lead block's weight pulls the ice lower, so more water gets displaced. If the lead fell out, the ice would pop up and displace less, just as in the spike case." Users of these stimuli may therefore wish to treat Bergs 2 similarly to Probability 1 and Switch 1 in that all contain fairly blatant logical errors.

C.11 Switch

Switch is also a variation on a Barnes et al. (2020)-investigated *Thinking Physics* question. Here, Dishonest's ball-on-a-string analogy is misleading. As Honest notes, the ball speeds up because the person pulling the string does work, tension has a component along the ball's inward displacement. On a fixed, frictionless track, no agent is pulling anything inward, so the mechanism that drives the speed-up doesn't exist. Dishonest asserts that 'the same principle applies' without explaining what that principle is or how it transfers from a string under tension to a fixed rail.

Dishonest also argues that because curvature changes along the connecting track, the contact force can't remain perpendicular to velocity. However, changing curvature doesn't change the contact force's orientation relative to velocity; it remains perpendicular at every point. Dishonest is left with little option but to hope that this is not self-evident to the judge, and to point out various irrelevant factors that change while on the connecting track without spelling out exactly how these lead to a change in speed. In Switch 1, Dishonest makes the blunder described at the end of Section 5.2; in addition, Honest directly disproves the claim that Dishonest introduces the formula to support (although judges may find the mathematics difficult to follow), and makes use of a visual analogy intended to make the perpendicularity of the contact force and the velocity more obvious. While Dishonest avoids these errors in Switch 2, they instead enter into a fairly direct contradiction by conceding a critical point about the axle geometry, which likely

[7] The original behavior in MALT-public 338159 was to overwrite the tests file in a more obvious way; for the purposes of the problem, the contents of what the file was overwritten with were modified to make the Dishonest position more plausible to participants without relevant domain knowledge.

explains the trend towards more correct judgments for this transcript.

C.12 Probability Puzzle

This problem was taken from Coelho (2024), which in turn is a restatement of a 'paradox' often attributed to Elchanan Mossel (Jin 2018). Many people find the correct answer unintuitive, because there is a simple, incorrect argument that can seem correct at first glance. In principle, Honest can fit the algebra of a partial proof of the correct answer into the debate limit, but in the first pilot this was found not to be an effective strategy, in part because judges felt unable to follow the details given the time provided. In the final transcripts used in the main study, in the half of each transcript consisting of the debate opened by Honest, Dishonest is forced to concede the main points of Honest's partial proof, but Dishonest doubles down on their original intuitive but incorrect argument and argues that the remainder of the proof does not go through. In the half containing the debate opened by Dishonest, Honest attempts to provide the intuition behind the correct answer; Dishonest is forced to resort to provably incorrect claims to respond, but the debate limit is reached before they can be fully disproved. Additionally, Probability Puzzle 1 was deliberately made easier by including an explicit contradiction in Dishonest's argument that Honest is able to point out before the end of the debate.

D Self-assessed knowledge questions

Given the topics of several of our questions, we administered a set of background questions and derived two binary indicators of whether participants potentially possessed some relevant (1) physics or (2) computer science or programming knowledge, to facilitate follow-up exploratory analysis. All participants were asked the following questions. Questions about the major and minor of a completed degree were shown only to participants who reported having completed an undergraduate, graduate, or doctorate degree; the question about a currently pursued major was shown only to participants who reported currently being in school.

1. **Highest level of education completed.** Options: Secondary education (e.g. GED/GCSE); High school/A-levels; Technical, vocational, or community college; Undergraduate degree (BA, BSc, or equivalent); Graduate degree (MA, MSc, MPhil, or equivalent); Doctorate degree (PhD, DPhil, or equivalent).
2. **Currently in school?** Options: Yes; No. (If Yes:) *What kind of education are you pursuing?* (same options as above).
3. **Primary subject (major) of highest completed degree.** (Shown to degree-holders.) Options: Physics; Computer Science / Software Engineering; Other (free response).
4. **Did you complete a minor in Physics?** (Shown to degree-holders.) Options: Yes; No.
5. **Primary subject (major) of the degree currently being pursued.** (Shown to those currently in school.) Options: Physics; Computer Science / Software Engineering; Other (free response).
6. **Have you ever taught physics?** Options: Yes, at the high-school level or above; Yes, but only below the high-school level; No.
7. **Do you consider yourself to be particularly knowledgeable about physics?** Options: Yes; No.
8. **Have you ever taught computer programming?** Options: Yes, at the high-school level or above; Yes, but only below the high-school level; No.
9. **Do you know how to program a computer?** Options: Yes; No.
10. **Have you ever spent 12 months or more employed as a software engineer or developer?** Options: Yes; No.

A participant was treated as having some physics knowledge if any of the following held:

- the major of their highest completed degree was Physics;
- the major of the degree they were currently pursuing was Physics;
- they reported completing a minor in Physics;
- they reported having taught physics at any level;
- they reported considering themselves particularly knowledgeable about physics.

A participant was treated as having some computer science or programming knowledge if any of the following held:

- the major of their highest completed degree was Computer Science / Software Engineering;
- the major of the degree they were currently pursuing was Computer Science / Software Engineering;
- they reported having taught computer programming at any level;
- they reported having spent 12 months or more employed as a software engineer or developer;
- they reported knowing how to program a computer.

E Diagnostics for the preregistered models

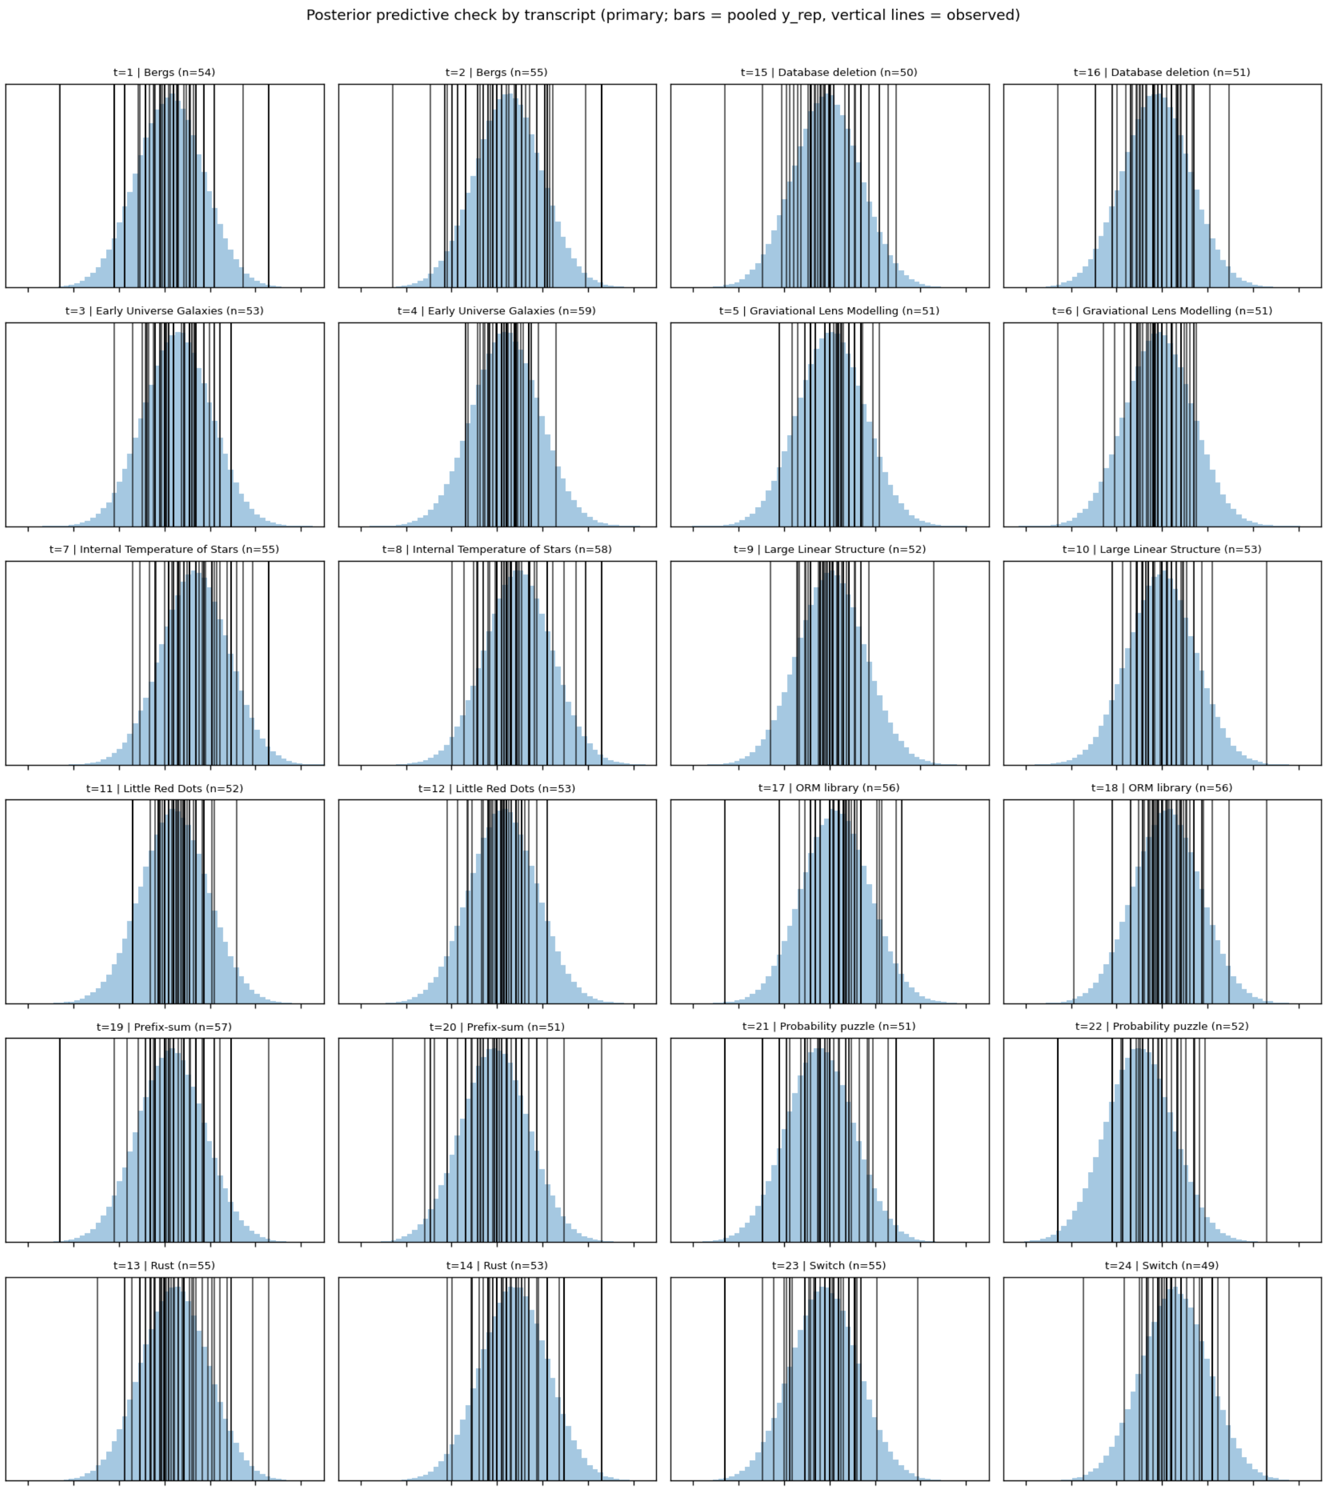

Figure 5: Posterior predictive check for the preregistered primary (Normal) model. Each panel shows pooled y_{rep} draws against the observed outcomes for one of the 24 transcripts. Outcome is the end-of-debate logit credence in the honest debater's position.

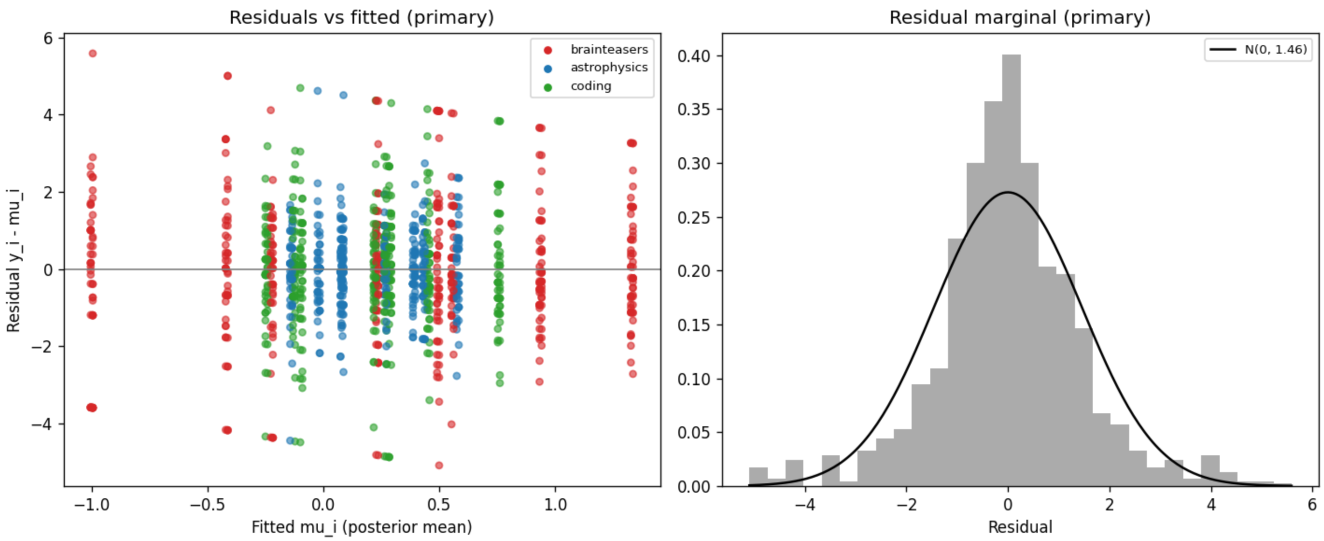

Figure 6: Residual diagnostics for the preregistered primary (Normal) model. **Left:** residuals $y_i - \hat{\mu}_i$ vs posterior-mean fitted values, coloured by domain. **Right:** marginal histogram of residuals overlaid with a Normal density whose SD equals the sample SD of the residuals.

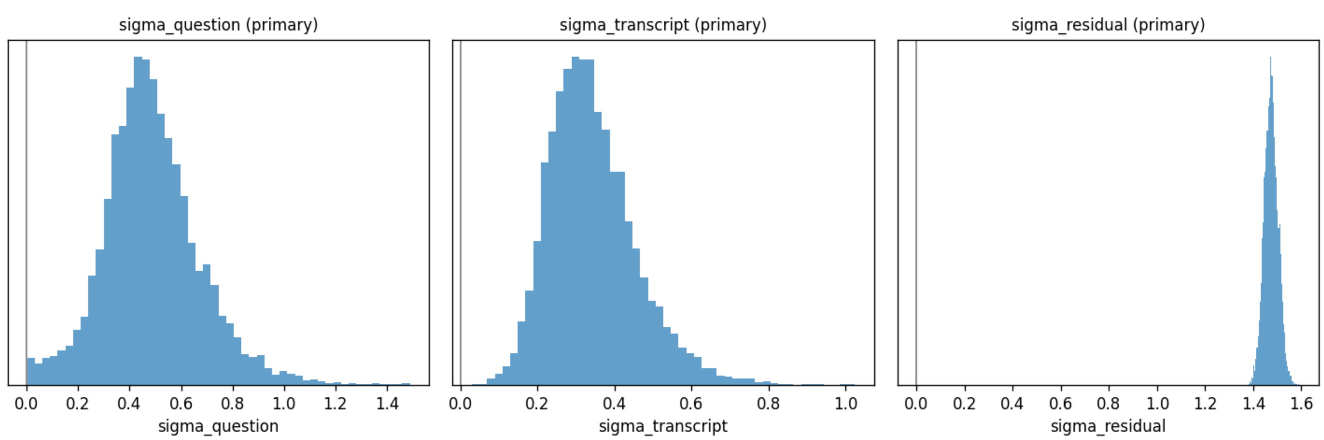

Figure 7: Variance-component posteriors for the preregistered primary (Normal) model. Panels show the posterior of σ_{question}, $\sigma_{\text{transcript}}$, and the residual scale σ.

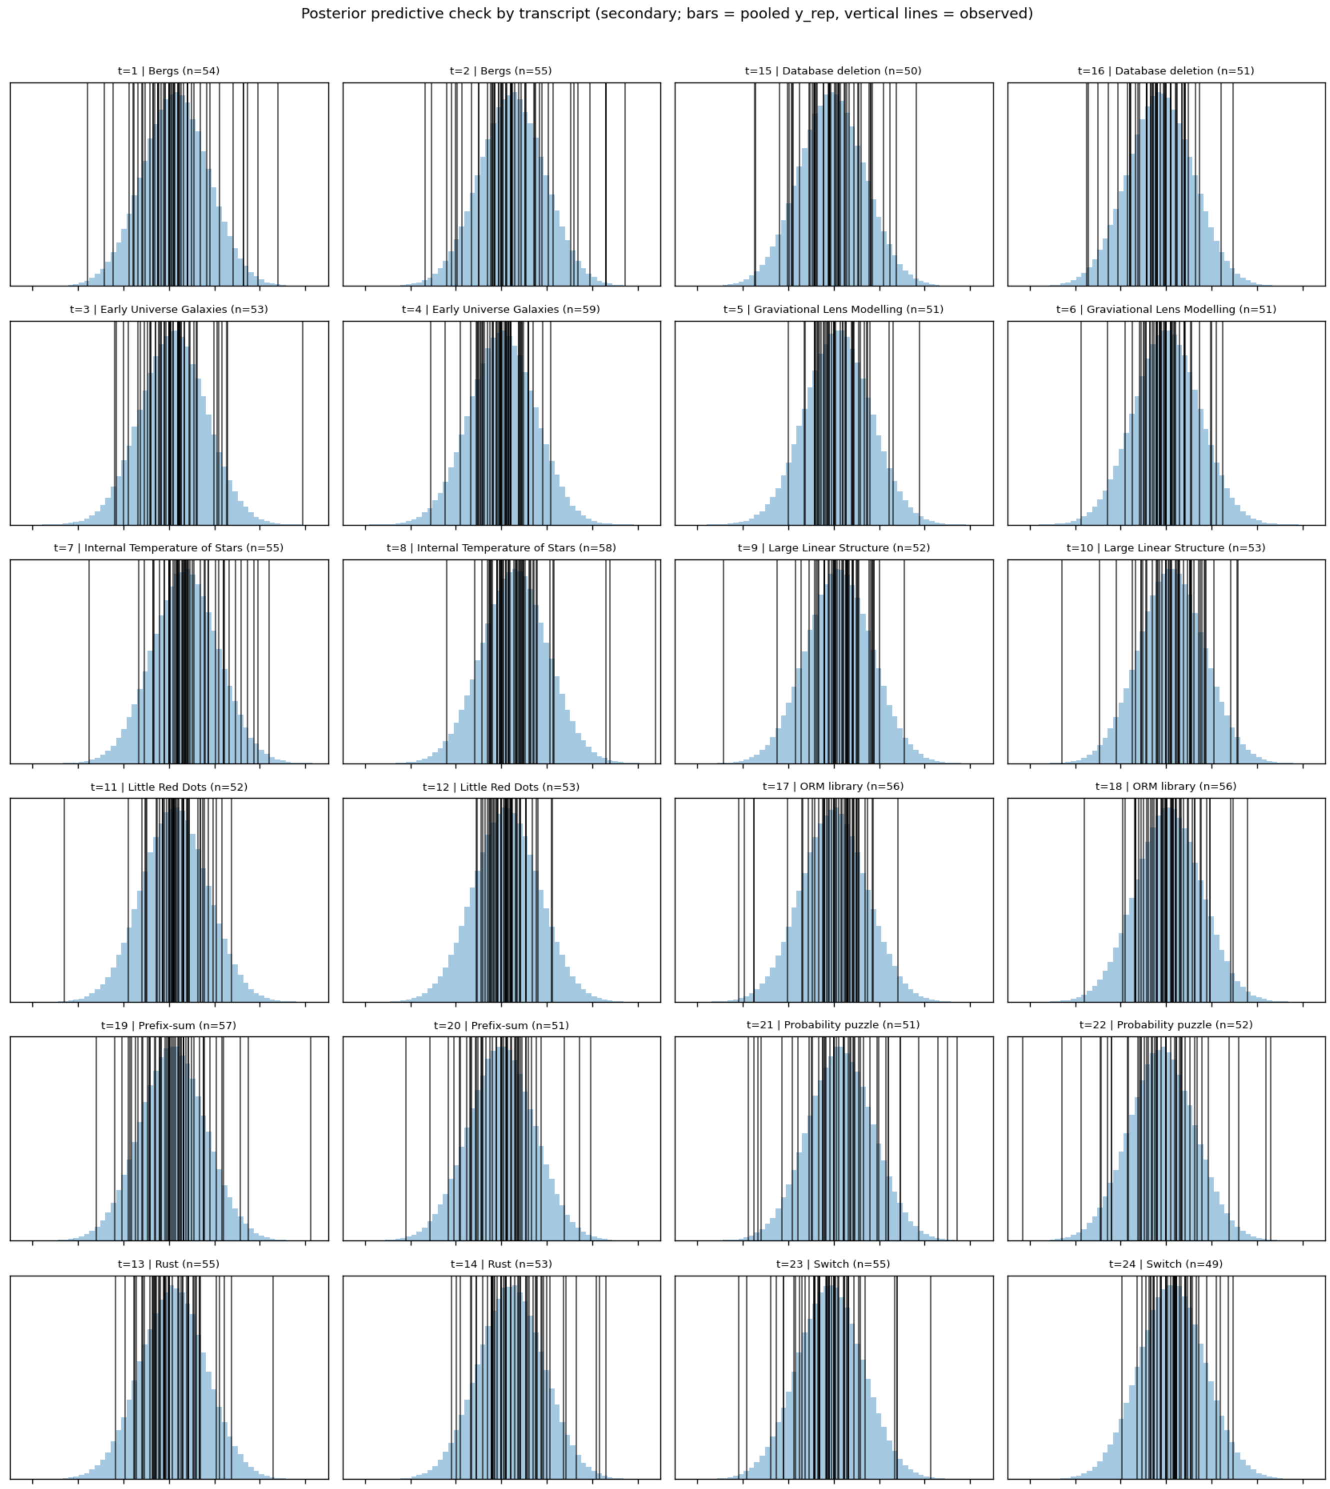

Figure 8: Posterior predictive check for the preregistered secondary (Normal) model. Outcome is the per-participant change in logit credence in the correct answer from section 1 to section 4.

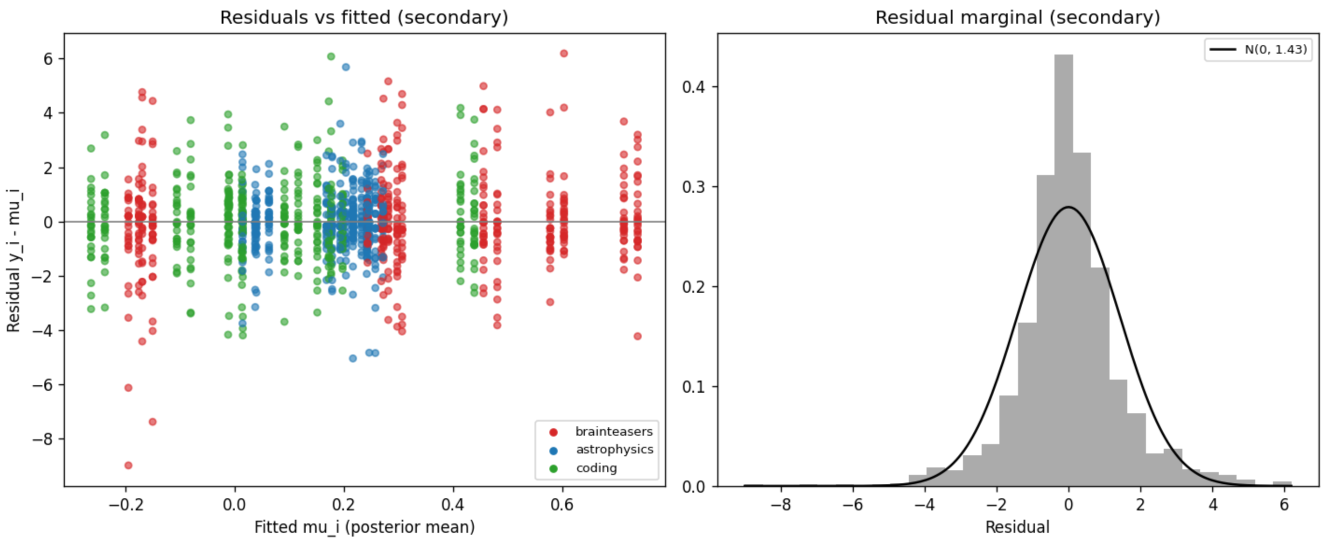

Figure 9: Residual diagnostics for the preregistered secondary (Normal) model. The marginal histogram (right) shows the heavy left tail that motivated the preregistered Student-t alternative reported in Appendix F.

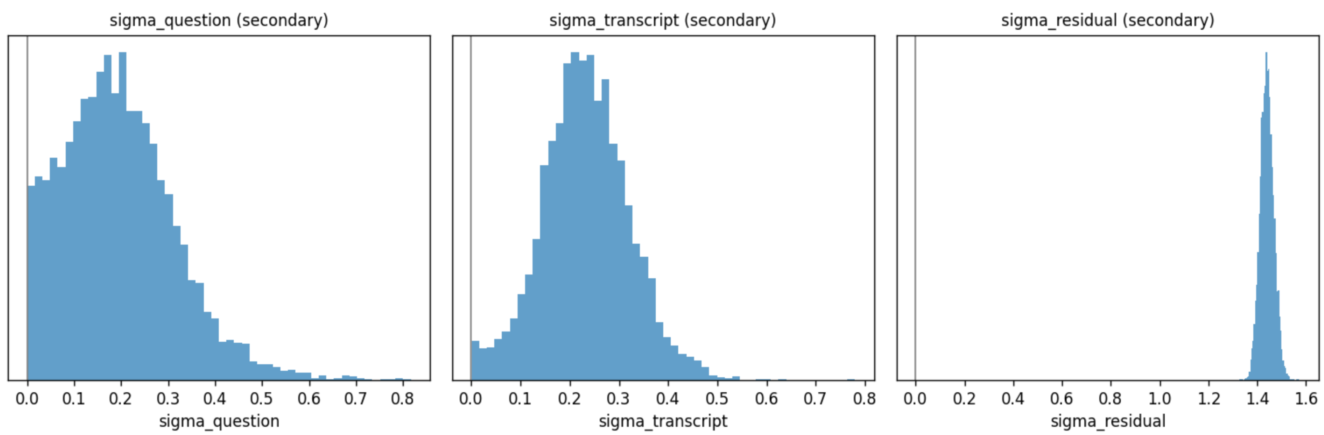

Figure 10: Variance-component posteriors for the preregistered secondary (Normal) model (σ_{question}, $\sigma_{\text{transcript}}$, σ_{residual}).

F Alternative models for the secondary outcome measure: diagnostics and model summaries

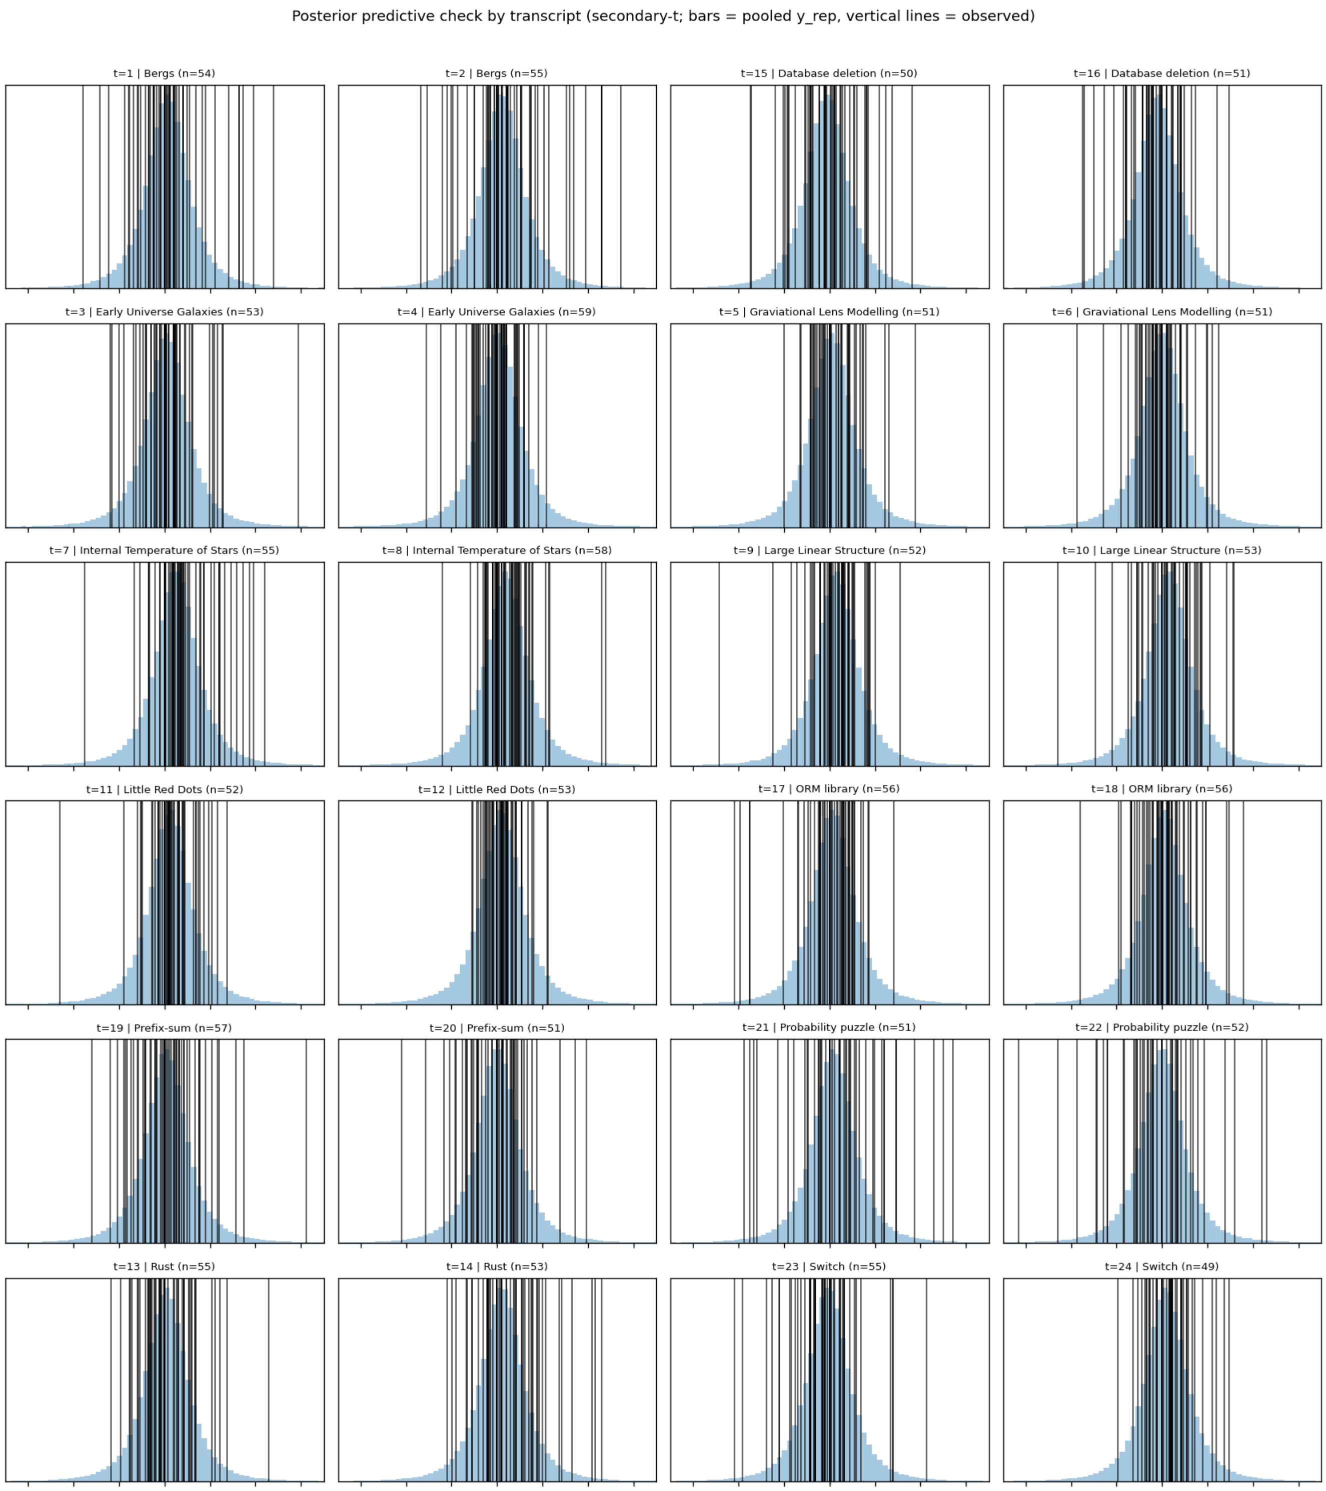

Figure 11: Posterior predictive check for the preregistered Student-t alternative to the secondary (Normal) model.

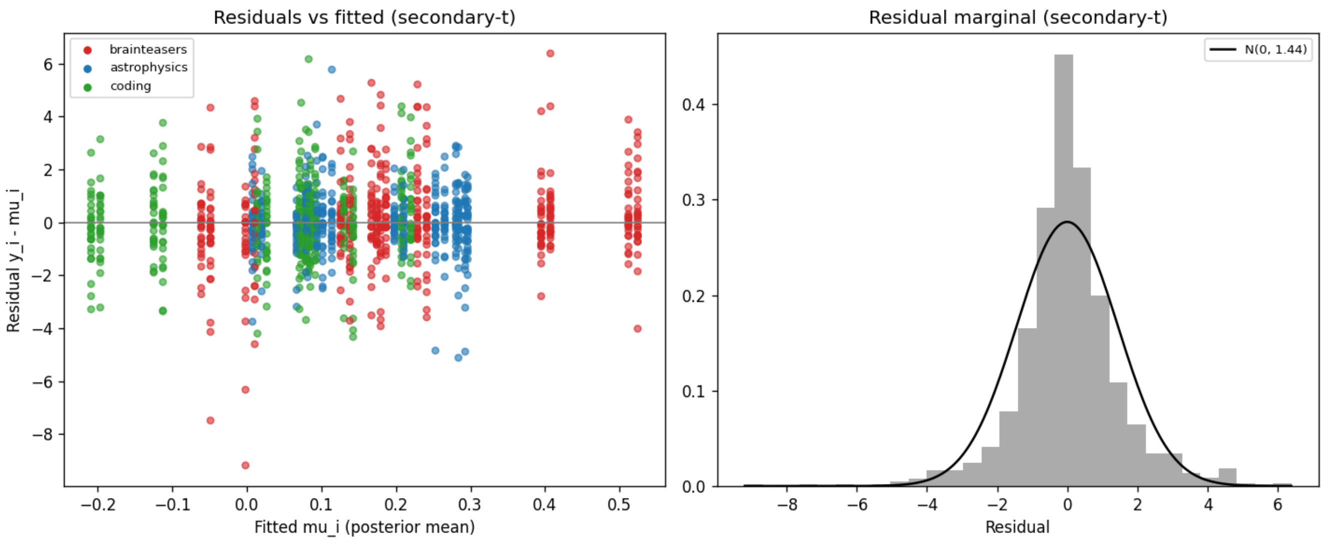

Figure 12: Residual diagnostics for the Student-t secondary model.

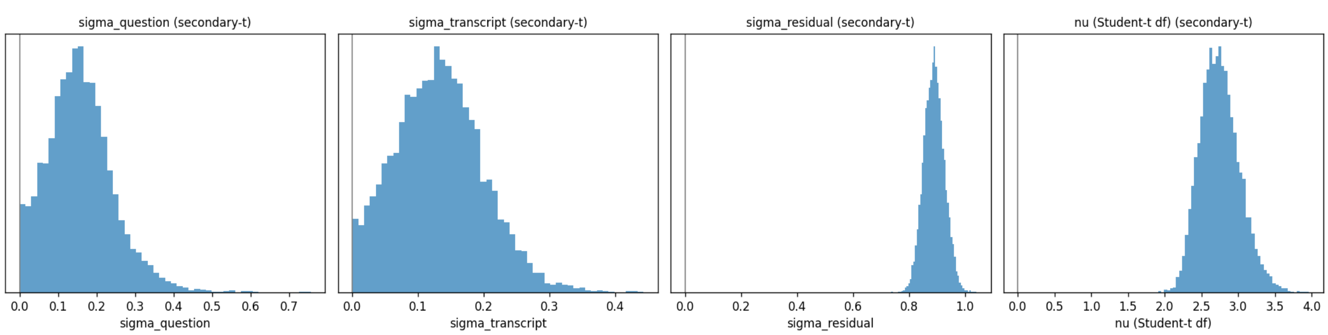

Figure 13: Variance-component posteriors for the Student-t secondary model, with the Student-t degrees of freedom ν shown as the rightmost panel. The posterior of $\sigma_{\text{transcript}}$ has shifted toward zero relative to the Normal fit.

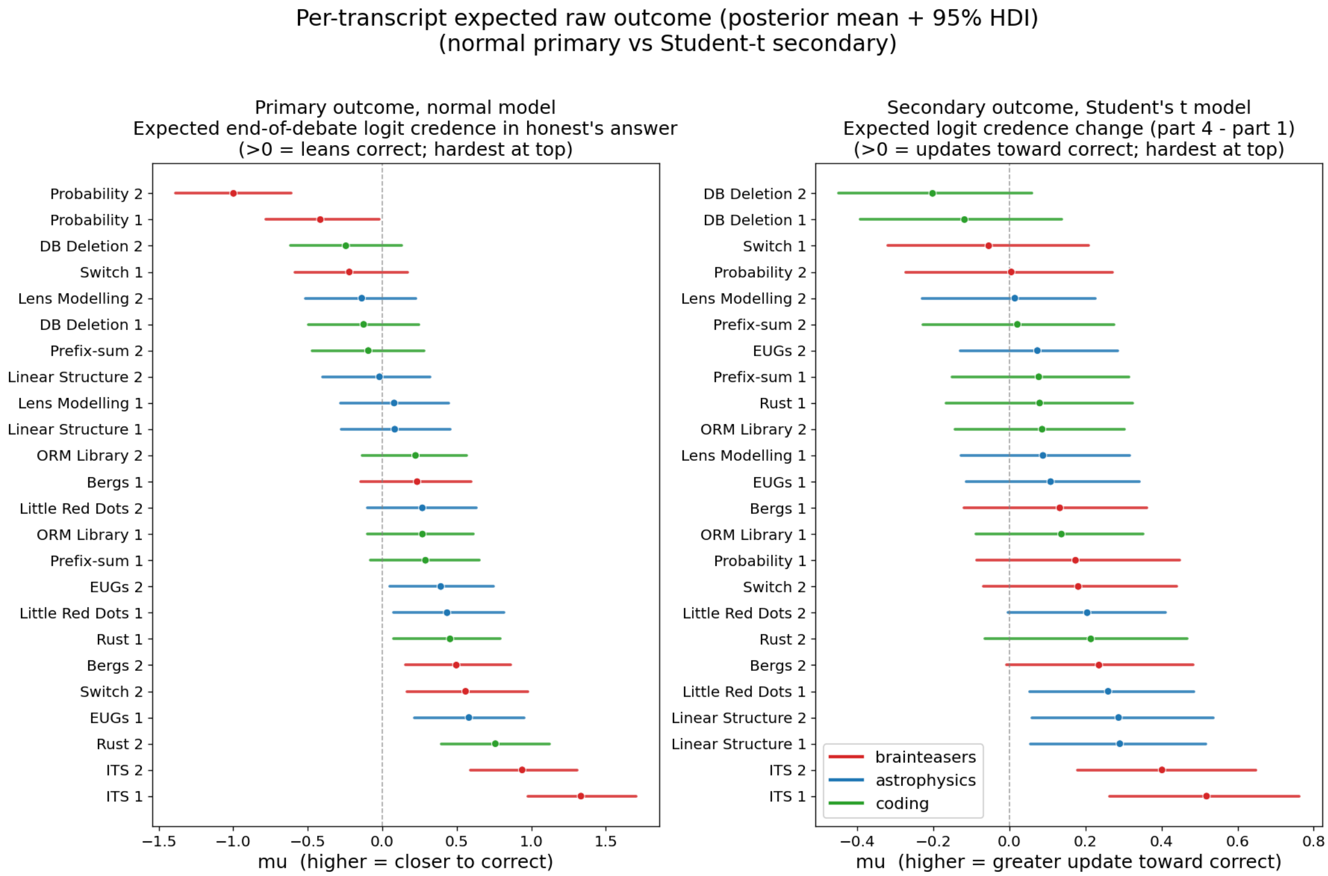

Figure 14: Per-transcript expected raw outcome (posterior mean and 95% HDI), comparing the preregistered primary fit to the Student-t secondary fit. **Left:** primary (Normal) outcome for reference. **Right:** secondary outcome under the Student-t alternative. Transcripts are ordered hardest-at-top within each panel.

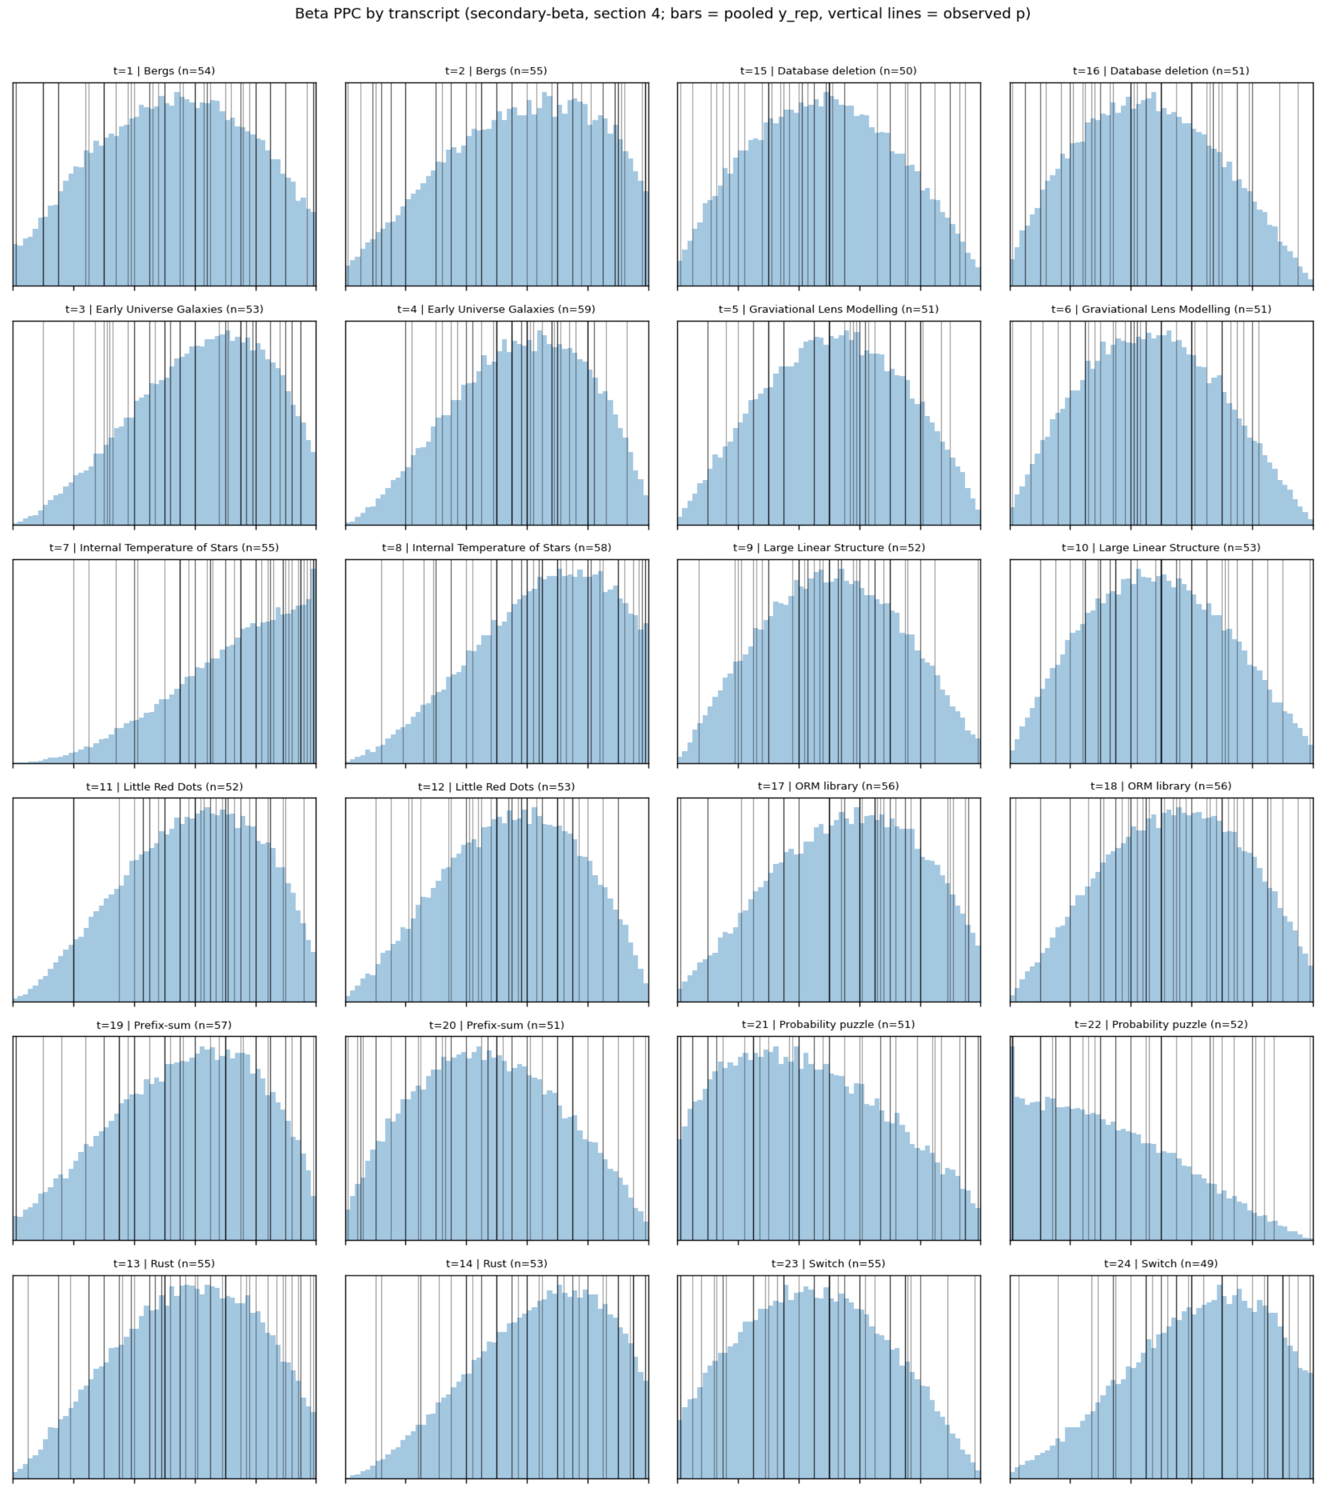

Figure 15: Posterior predictive check for the exploratory Beta secondary model, section 4. The Beta model is fit on the raw per-section credences in $[0.01, 0.99]$ rather than the clamped logit-credence change, eliminating the boundary pile-up that contributes to the Normal fit's heavy tail.

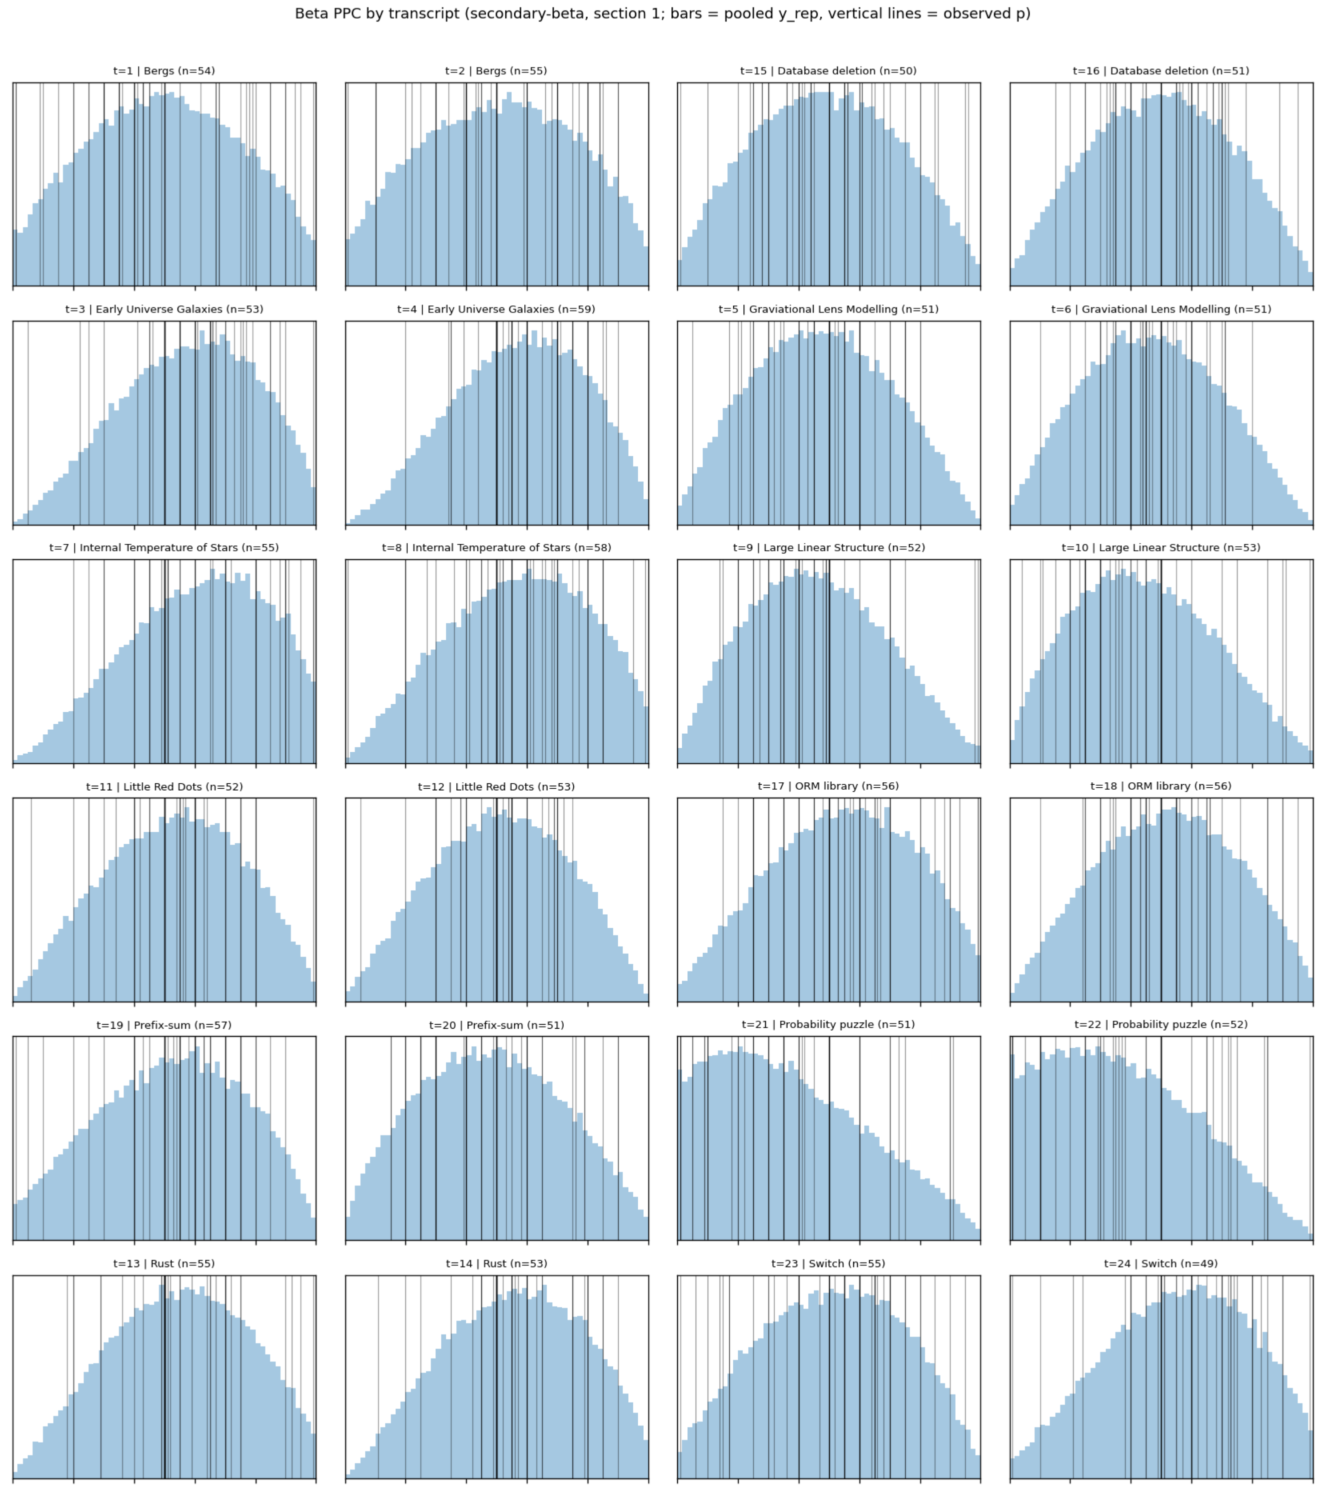

Figure 16: Posterior predictive check for the exploratory Beta secondary model, section 1.

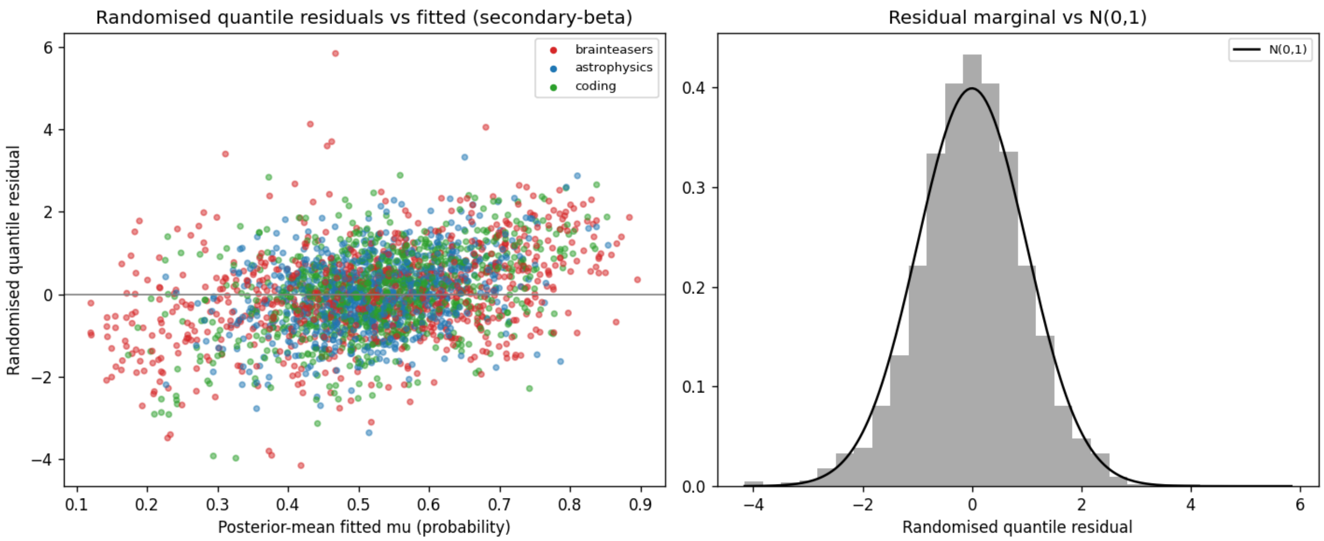

Figure 17: Residual diagnostics for the exploratory Beta secondary model. Because Beta residuals on the $(0, 1)$ scale are not directly comparable to the Normal / Student-t residuals, these are randomised quantile residuals, computed here by transforming each observation through its fitted Beta CDF at a randomly chosen posterior draw and then through Φ^{-1}.

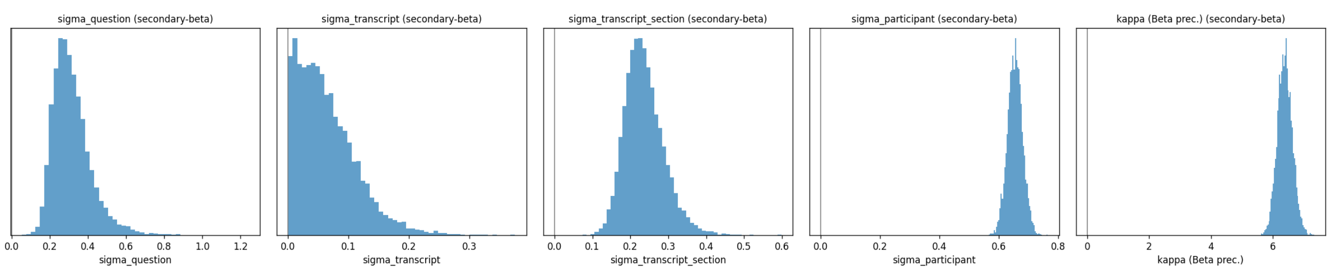

Figure 18: Variance-component posteriors for the exploratory Beta secondary model. The rightmost panel shows the Beta precision parameter κ. In Panel 3 we see the Beta fit recovers the between-transcript variance that the Student-t fit had compressed against zero, with $\sigma_{v,\text{sec}} = 0.237$ (95% HDI $[0.142, 0.341]$, $P(\sigma_{v,\text{sec}} < 0.1) \approx 0.001$).

F.1 Student-t secondary model: parameter summaries

Table 4: Posterior summaries for the preregistered Student-t alternative to the secondary (Normal) outcome model. Fixed effects and the residual scale are on the same logit-credence-change scale as the secondary Normal fit; ν is the Student-t degrees of freedom.

Parameter	Posterior mean	95% HDI
Intercept (α)	$+0.188$	$[-0.049, +0.423]$
order (B-honest vs. A-honest)	$+0.012$	$[-0.107, +0.126]$
domain: astrophysics	-0.030	$[-0.350, +0.285]$
domain: coding	-0.156	$[-0.470, +0.161]$
σ_q (question)	0.158	$[\ \ 0.000, +0.311]$
σ_v (transcript)	0.132	$[\ \ 0.000, +0.246]$
σ_{resid}	0.890	$[\ \ 0.820, +0.962]$
ν (Student-t df)	2.754	$[\ \ 2.263, +3.302]$

F.2 Beta secondary model: parameter summaries

Table 5: Posterior summaries for the exploratory Beta alternative to the secondary outcome model. Fixed effects and variance components are on the Beta-logit scale (the link-function scale for the Beta likelihood). κ is the Beta precision parameter, with larger κ implying tighter concentration of p around its fitted mean.

Parameter	Posterior mean	95% HDI
Intercept (α)	-0.007	$[-0.360, +0.319]$
order (B-honest vs. A-honest)	$+0.002$	$[-0.092, +0.093]$
β_{section} (mean §1→§4 update)	$+0.140$	$[+0.025, +0.254]$
domain: astrophysics	$+0.032$	$[-0.443, +0.513]$
domain: coding	$+0.092$	$[-0.381, +0.577]$
σ_q (question)	0.316	$[\ \ 0.150, +0.520]$
σ_v (transcript intercept)	0.064	$[\ \ 0.000, +0.156]$
$\sigma_{v,\text{sec}}$ (transcript × §4 slope)	0.237	$[\ \ 0.142, +0.341]$
σ_p (participant)	0.655	$[\ \ 0.606, +0.705]$
κ (Beta precision)	6.383	$[\ \ 5.910, +6.849]$

G Energy-transition diagnostics

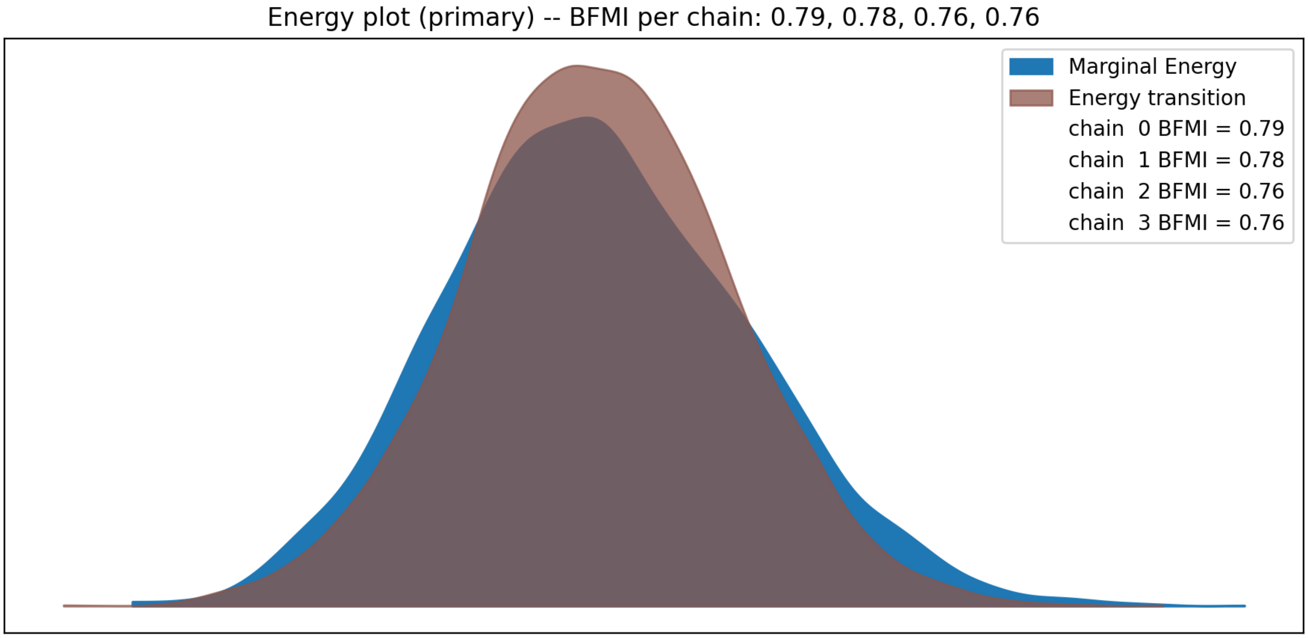

Figure 19: Energy-transition plot for the preregistered primary (Normal) model. The marginal energy distribution π_E is overlaid with the distribution of energy transitions $\pi_{\Delta E}$ in each chain. Close agreement between the two indicates that NUTS is exploring efficiently, consistent with the per-chain E-BFMI values reported in the main text.

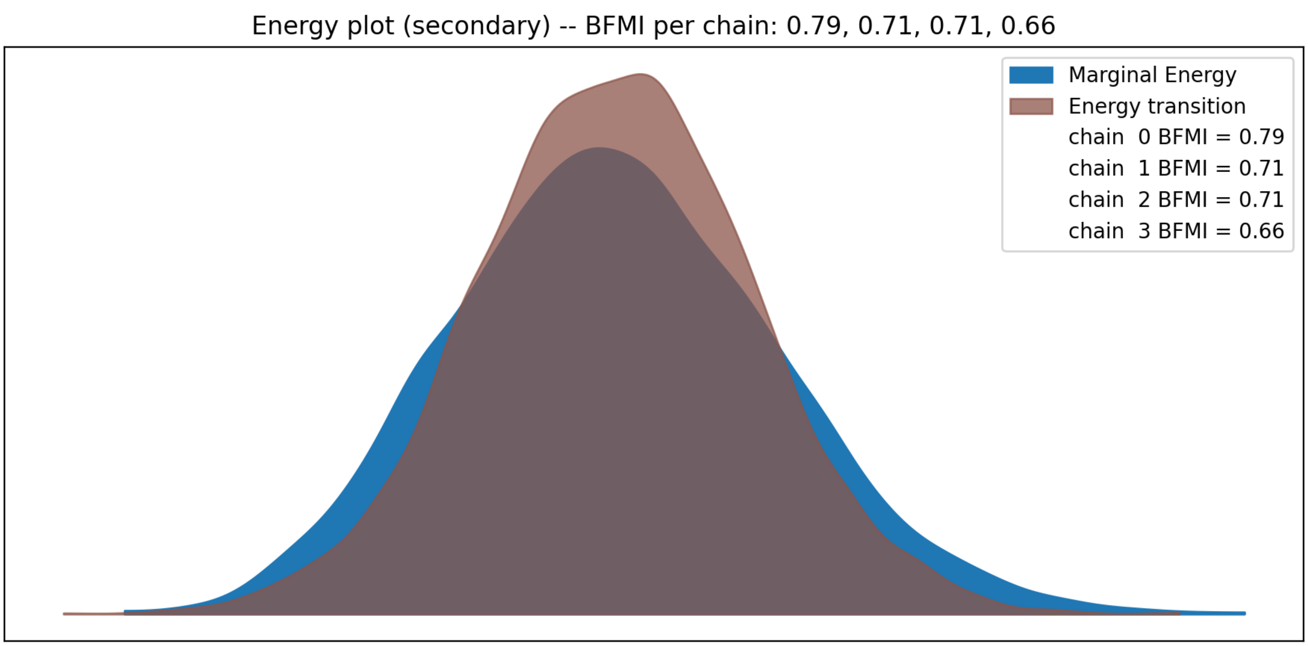

Figure 20: Energy-transition plot for the preregistered secondary (Normal) model.

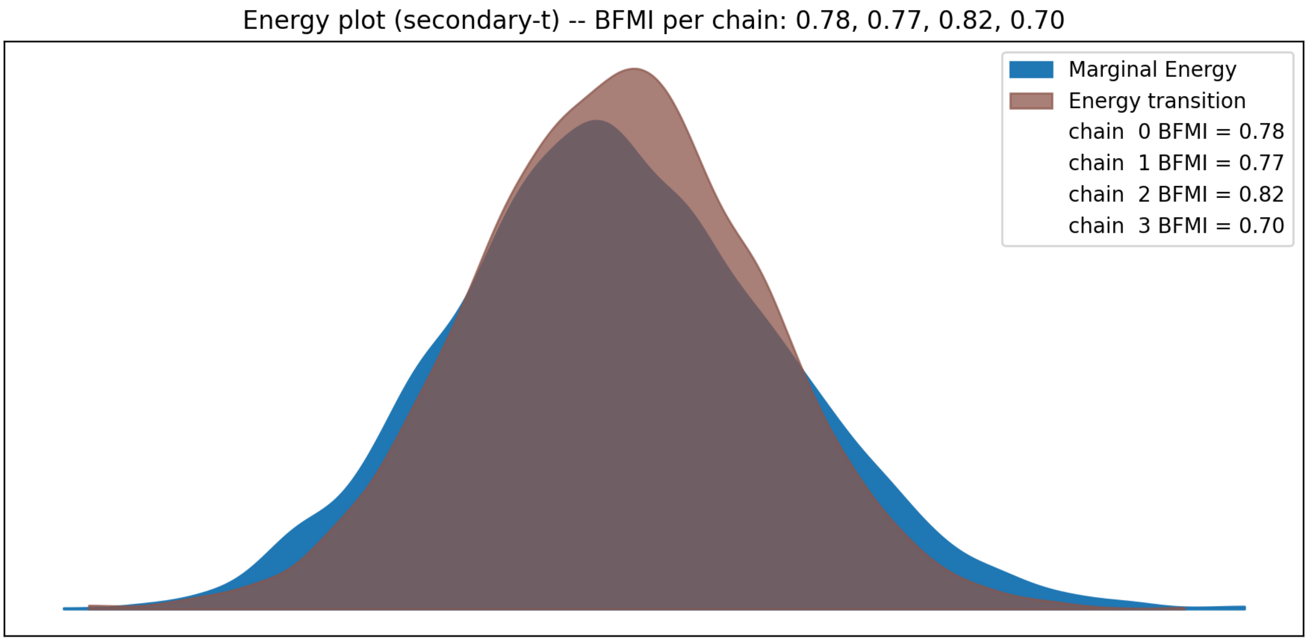

Figure 21: Energy-transition plot for the preregistered Student-t alternative to the secondary (Normal) model.

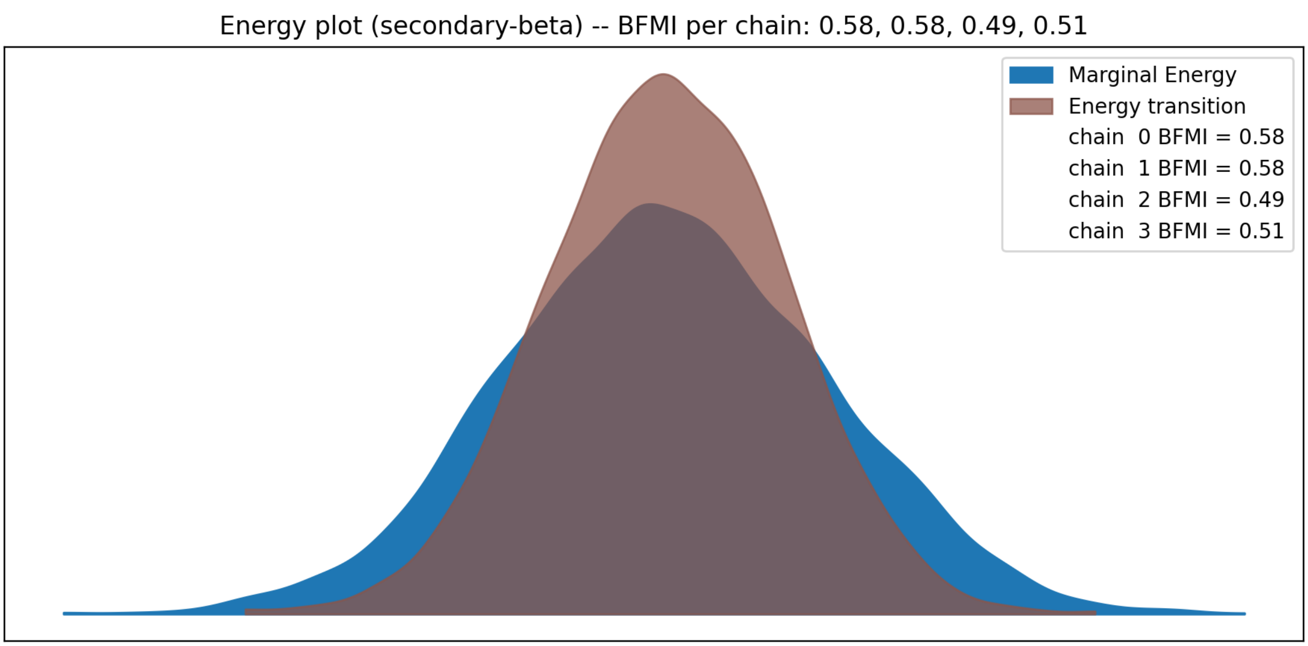

Figure 22: Energy-transition plot for the exploratory Beta secondary model.

H Alternative tier assignments and tier-membership uncertainty

Modal Tier × Domain, secondary outcome model (preregistered normal)

	Tier 1 (easiest)	Tier 2	Tier 3	Tier 4 (hardest)
Astrophysics		Little Red Dots 1 Large Linear Structure 2 Little Red Dots 2 Large Linear Structure 1 Early Universe Galaxies 1 Gravitational Lens Modelling 1	Early Universe Galaxies 2 Gravitational Lens Modelling 2	
Brainteasers	Internal Temperature of Stars 1 Internal Temperature of Stars 2 Bergs 2	Probability Puzzle 1 Bergs 1 Switch 2		Switch 1 Probability Puzzle 2
Reward Hacking	Rust 2	Rust 1	Prefix-sum 1 ORM Library 2	ORM Library 1 Prefix-sum 2 Database Deletion 1 Database Deletion 2

Figure 23: Modal difficulty-tier assignment for each of the 24 transcripts under the preregistered secondary (Normal) outcome model, cross-tabulated by domain. Tier 1 contains transcripts on which judges updated most toward the correct answer over the course of the debate; Tier 4 contains those on which judges updated least (or away from the correct answer). Layout matches Figure 1 in the main text.

Modal Tier × Domain, secondary outcome model (preregistered Student's t)

	Tier 1 (easiest)	Tier 2	Tier 3	Tier 4 (hardest)
Astrophysics	Large Linear Structure 1 Large Linear Structure 2 Little Red Dots 1	Little Red Dots 2	Early Universe Galaxies 1 Gravitational Lens Modelling 1 Early Universe Galaxies 2	Gravitational Lens Modelling 2
Brainteasers	Internal Temperature of Stars 1 Internal Temperature of Stars 2 Bergs 2	Switch 2 Probability Puzzle 1 Bergs 1		Probability Puzzle 2 Switch 1
Reward Hacking		Rust 2 ORM Library 1	ORM Library 2 Rust 1 Prefix-sum 1	Prefix-sum 2 Database Deletion 1 Database Deletion 2

Figure 24: Modal difficulty-tier assignment under the Student-t secondary model. Layout as in Figure 23.

Modal Tier × Domain, secondary outcome model (exploratory beta)

	Tier 1 (easiest)	Tier 2	Tier 3	Tier 4 (hardest)
Astrophysics		Little Red Dots 1 Large Linear Structure 1 Early Universe Galaxies 1 Gravitational Lens Modelling 1 Large Linear Structure 2	Little Red Dots 2 Early Universe Galaxies 2	Gravitational Lens Modelling 2
Brainteasers	Internal Temperature of Stars 1 Internal Temperature of Stars 2 Bergs 2 Switch 2	Probability Puzzle 1	Bergs 1	Probability Puzzle 2 Switch 1
Reward Hacking	Rust 2	Prefix-sum 1	Rust 1 ORM Library 2 ORM Library 1	Database Deletion 1 Prefix-sum 2 Database Deletion 2

Figure 25: Modal difficulty-tier assignment under the exploratory Beta secondary model. Layout as in Figure 23.

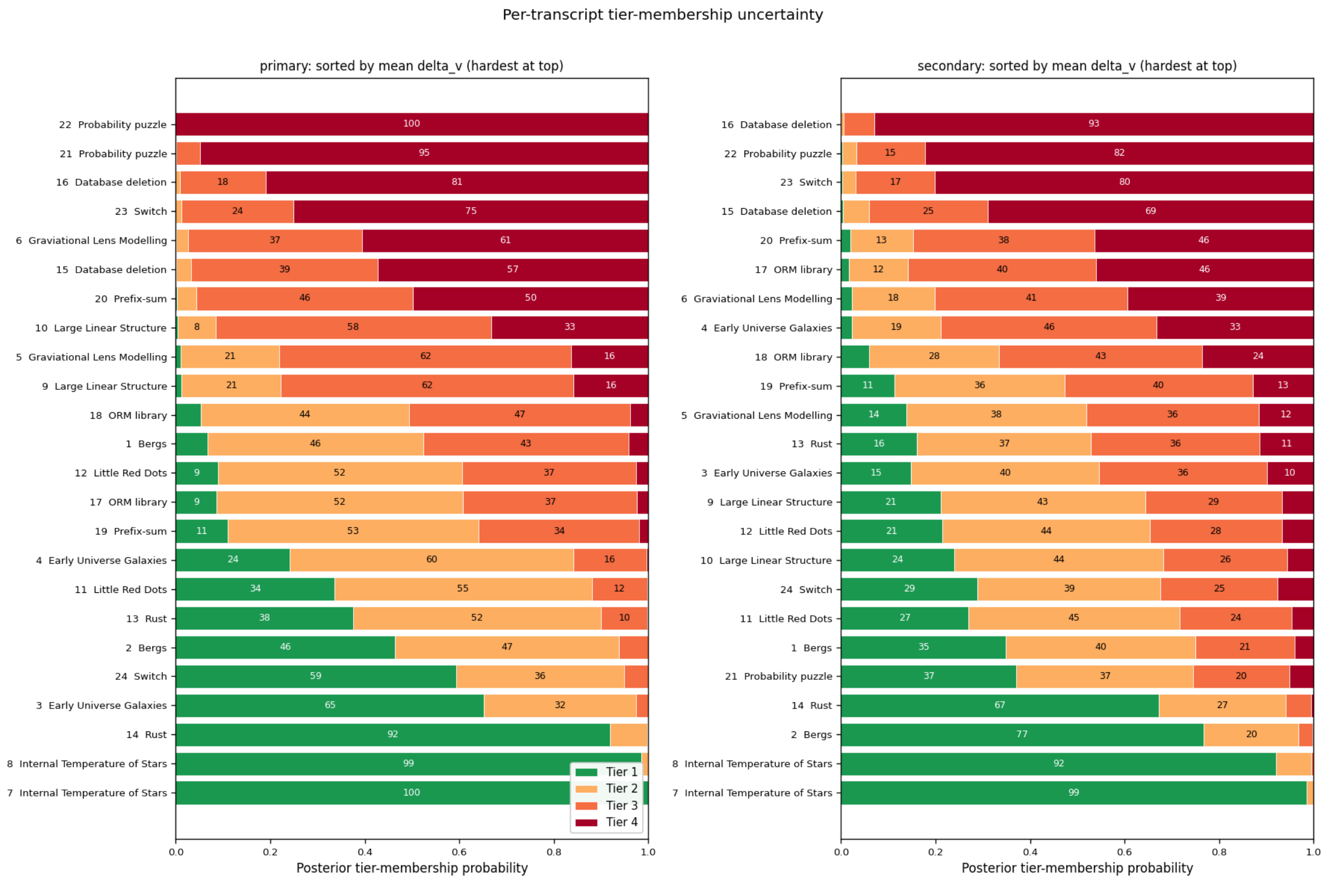

Figure 26: Per-transcript posterior tier-membership probabilities under the preregistered primary (Normal; **left**) and secondary (Normal; **right**) outcome models. Each row is one of the 24 transcripts; the stacked horizontal bar shows the posterior probability of that transcript falling in Tier 1 (green, easiest) through Tier 4 (red, hardest).

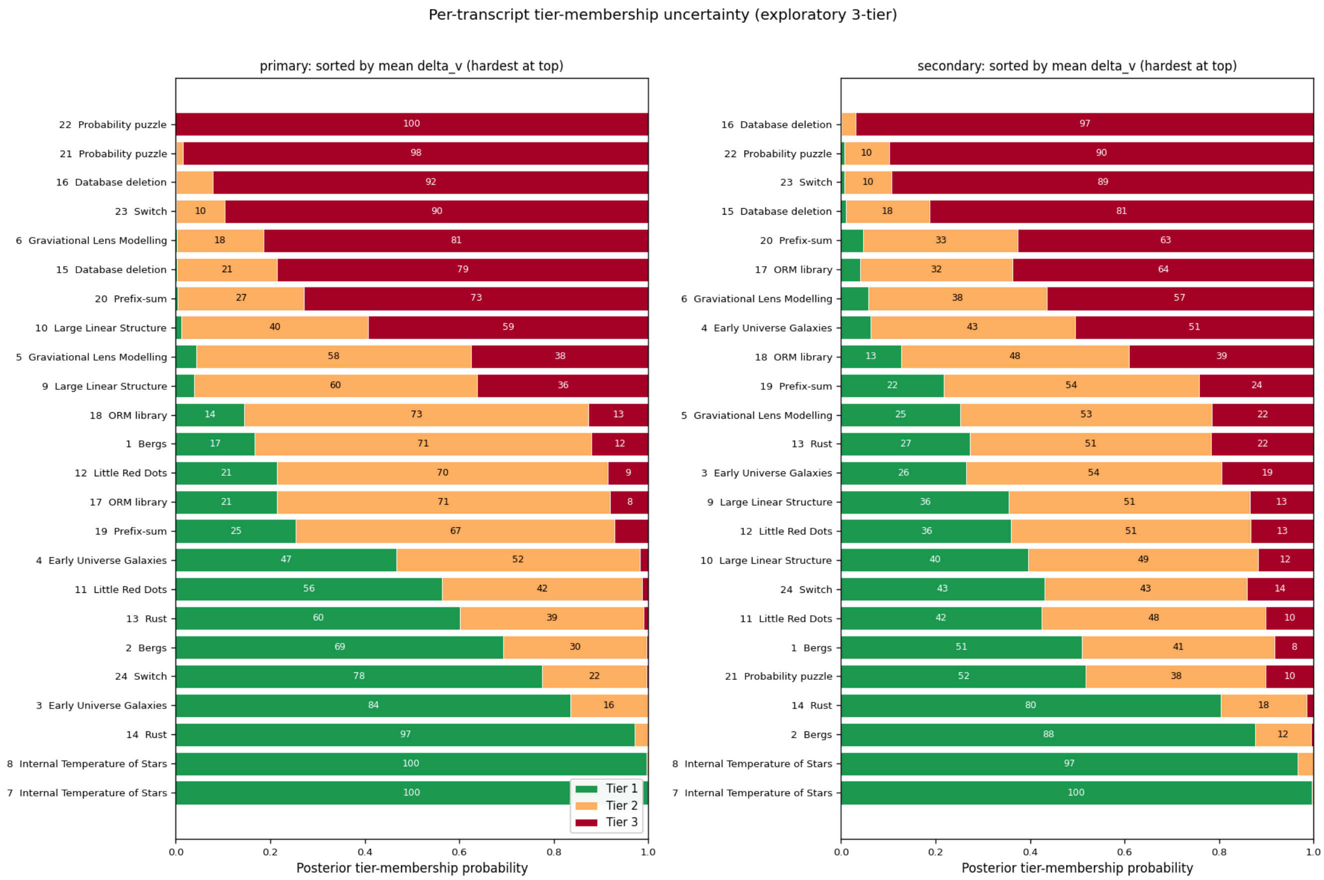

Figure 27: Per-transcript posterior tier-membership probabilities for the preregistered primary (**left**) and secondary (**right**) Normal models under the exploratory three-tier classification (Tier 1 easiest, Tier 3 hardest).

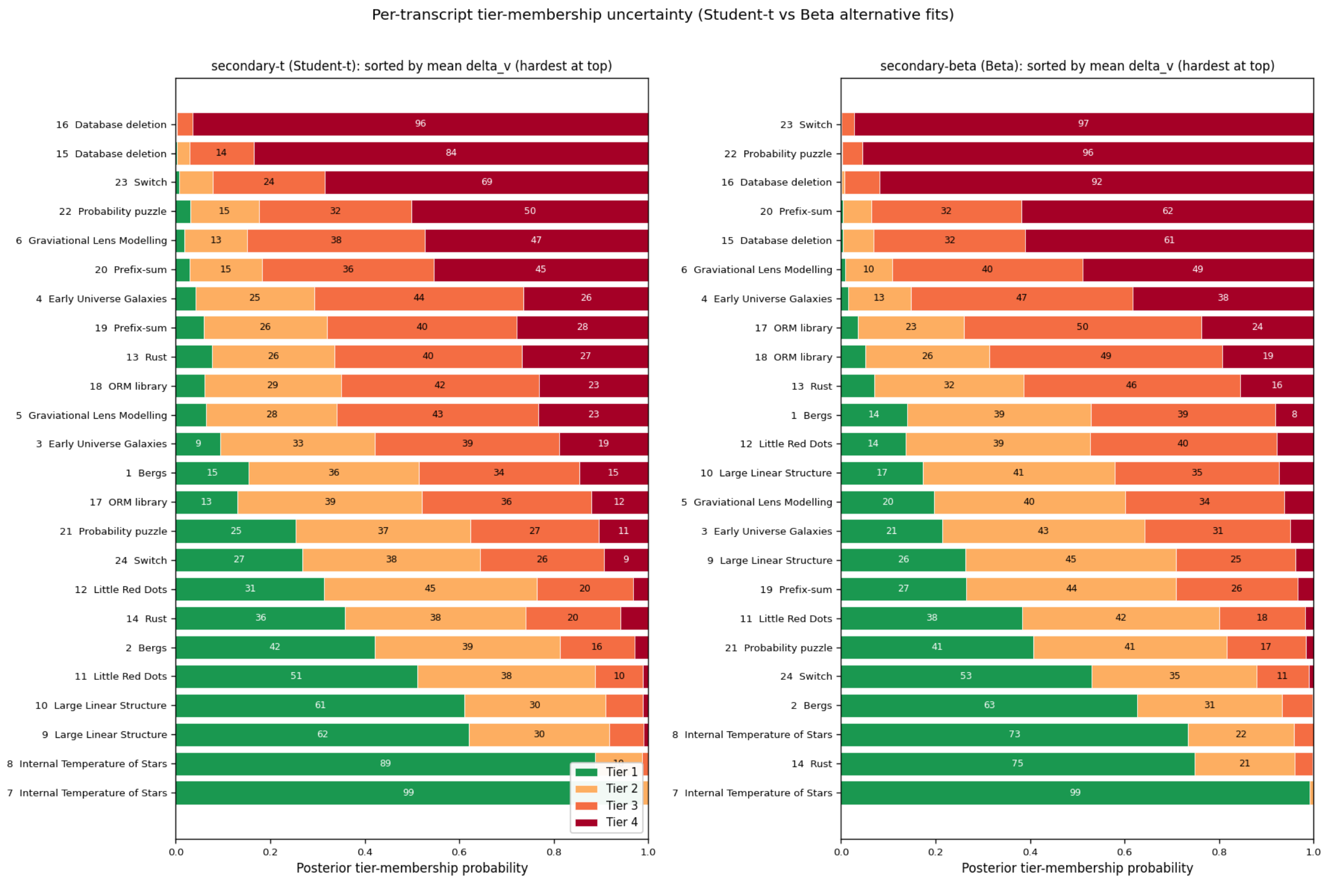

Figure 28: Per-transcript posterior tier-membership probabilities under the two alternative fits for the secondary outcome measure: Student-t (**left**) and exploratory Beta (**right**).

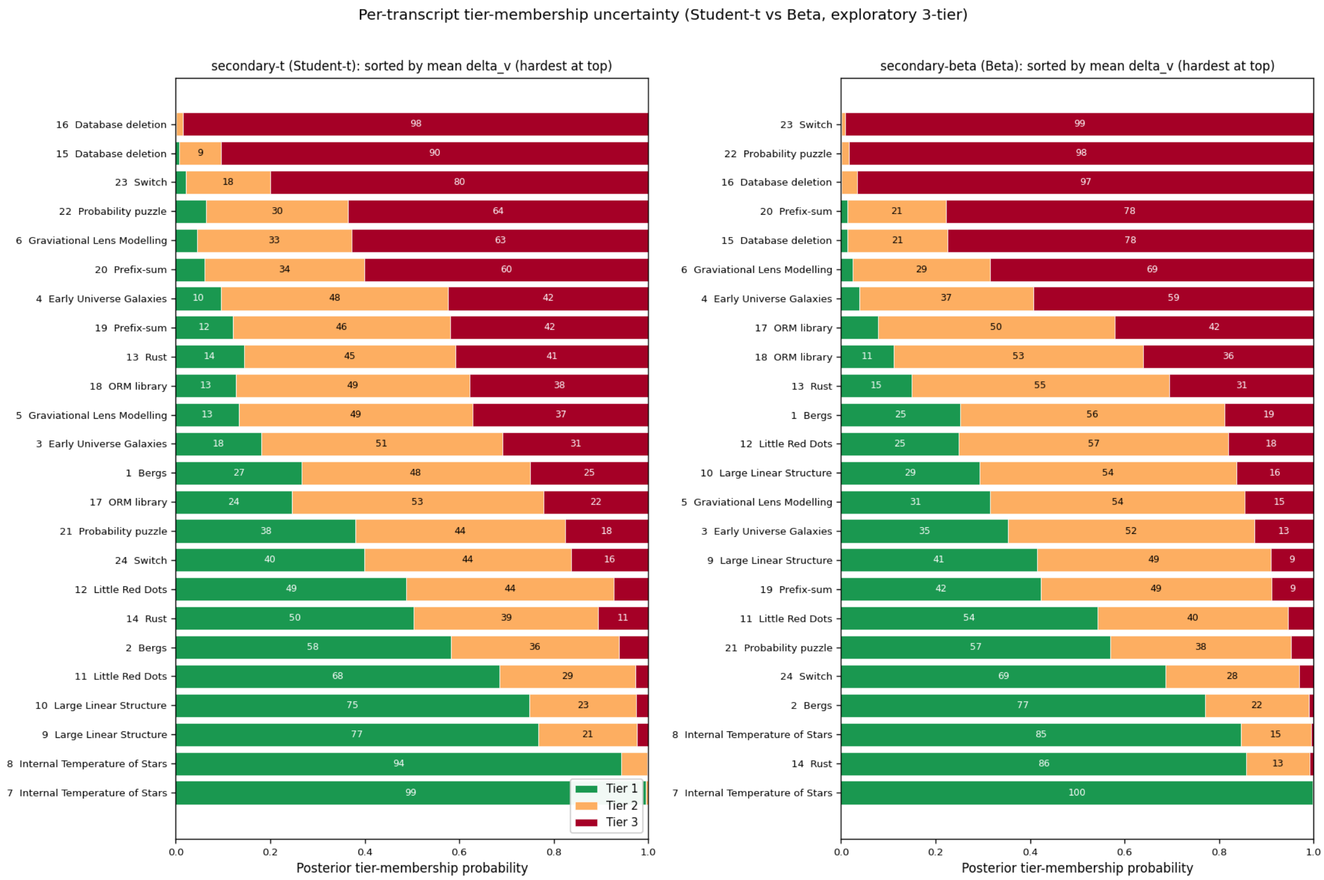

Figure 29: Per-transcript posterior tier-membership probabilities under the Student-t (**left**) and Beta (**right**) alternative secondary fits, under the exploratory three-tier classification.